

HVK: A Tensor-Network and Quantum-Observable Architecture for Spatially Aware Image Reconstruction

Sparsho Chakraborty* and Siddhartha Patra†

Abstract—We present the Hamiltonian Vision Kernel (HVK), a hybrid image reconstruction framework that combines matrix-product-state (MPS) patch features, positional encodings, parameterized variational quantum circuits, Pauli-observable latent vectors, and Hamiltonian energy diagnostics. The architecture was designed to test whether a quantum-observable bottleneck can provide useful patch-specific structure beyond a classical decoder alone. We report a controlled ablation study on the *Monalisa* benchmark covering frozen classical/quantum parameters, random and classical VQC replacements, no-entanglement circuits, Hamiltonian-loss removal, observable-noise removal, ZZ-only observables, qubit-count sweeps, MPS bond-dimension sweeps, and single-image versus multi-image generalization. The results show that the observable channel is load-bearing: shuffling or zeroing observables sharply degrades reconstruction. However, the same ablations do not support a claim of quantum advantage in the current implementation: fixed quantum features, non-entangling circuits, and parameter-matched classical replacements reconstruct at or above the trained VQC baseline, while removing the Hamiltonian energy term slightly improves PSNR. HVK is therefore best characterized as an experimentally transparent observable-latent reconstruction platform, with the Hamiltonian currently serving as an interpretable diagnostic rather than the dominant source of reconstruction accuracy.

Index Terms—Quantum Image Processing, Tensor Networks, Variational Quantum Algorithms, Pauli Observables, Hamiltonian Diagnostics, Image Reconstruction.

I. INTRODUCTION

Classical approaches—such as regularized inversion and, more recently, deep convolutional priors—have driven decades of progress, but their computational and memory costs scale poorly with image resolution and dimensionality. This has motivated a continuing search for representations and algorithms with more favorable scaling.

Quantum computing has emerged as one candidate paradigm. Variational quantum algorithms have demonstrated that parameterized quantum circuits, optimized in a hybrid quantum–classical loop, can prepare physically meaningful states on near-term hardware [1], [2]. The variational quantum eigensolver and its descendants—including the variational Hamiltonian ansatz and ADAPT-VQE—exploit the local decomposability of physically relevant Hamiltonians to construct shallow, problem-aware circuits whose parameters have transparent interpretations as rotation angles in a structured ansatz

[1]–[3]. These developments raise a natural question: can Hamiltonian-guided variational principles, originally designed for electronic-structure problems, be repurposed for image reconstruction, where the “Hamiltonian” encodes imaging physics rather than electronic interactions?

Quantum image processing (QIMP) provides the encoder side of such a pipeline. The flexible representation of quantum images encodes pixel colors in qubit rotation angles and pixel positions in computational-basis states, while the novel enhanced quantum representation (NEQR) stores grayscale values deterministically in basis states using q -qubit registers [4]. Successive works have improved and extended these encodings: quantum realizations of classical interpolation have been demonstrated for NEQR [5], and full quantum edge-detection pipelines for NEQR-encoded images have been validated theoretically and experimentally [6]. Quantum probability image encodings further reduce qubit and gate costs while preserving the structure needed for downstream processing [6]. However, most QIMP pipelines treat the encoded image as an opaque state to be acted upon; they do not expose the spatial correlations among neighboring pixels in a form that admits efficient variational optimization.

Tensor networks provide exactly this missing structure. Originally developed as compact representations of many-body quantum states, matrix product states (MPS), projected entangled pair states (PEPS), and the multi-scale entanglement renormalization ansatz (MERA) have proven remarkably effective as compressed representations of structured classical data, including images [7]. Classical tensor-network methods have been applied to image compression [8], to probabilistic generative modeling under partial observation [9], and—relevantly for inverse problems—to compressed sensing reconstructions whose latent state is itself a tensor network [10]. MPS/PEPS-structured quantum circuits have also been proposed as classical–quantum bridges that are qubit-efficient and well-suited to circuit-cutting techniques on near-term devices [11].

Despite this rich prior work, two streams remain largely disconnected. On one side, classical tensor-network formulations for imaging encode the image as an MPS and perform contraction and optimization without invoking quantum Hamiltonian dynamics; on the other, variational quantum algorithms for image reconstruction encode the image as a generic parameterized state without exposing the spatial bond structure of a tensor network [12]–[14]. The intersection—a Hamiltonian whose evolution respects and is explicitly coupled to a tensor-

*School of Basic Sciences, Department of Physics, Indian Institute of Technology Bhubaneswar, Odisha, India.

† Assistant Professor, Centre for Quantum Engineering, Research and Education (CQUERE), TCG CREST, Kolkata, India.

network topology that mirrors pixel adjacency—has not been systematically developed.

In this paper we introduce the HVK architecture, a Hamiltonian-guided tensor-network framework for spatially aware quantum image reconstruction. Rather than treating the image-encoding circuit and the reconstruction prior as separate modules, HVK fuses them: a problem-specific Hamiltonian is constructed such that its low-energy subspace, when probed via a variational ansatz organized as a tensor network, preferentially contains states consistent with both the measurement physics and the spatial structure of the scene.

II. THEORETICAL FRAMEWORK

A. Hamiltonian Operator Dynamics

HVK treats each spatial patch of an input image as a quantum many-body state on a finite spin-1/2 chain, encodes that state as a matrix product state, and exposes a Hamiltonian whose ground-state structure induces spatially aware correlations among neighboring pixels. This establishes the operator-theoretic foundations: the Heisenberg family of Hamiltonians, the corresponding energy objective, and the symmetry structure inherited by the resulting observable manifold.

1) *The Heisenberg Family of Hamiltonians:* For a one-dimensional chain of n_{bonds} bonds indexed by i and sites $(i, i+1)$, the anisotropic Heisenberg Hamiltonian is defined as:

$$H_{\text{Heis}} = \sum_{i=0}^{n_{\text{bonds}}-1} \left(J_x^{(i)} X_i X_{i+1} + J_y^{(i)} Y_i Y_{i+1} + J_z^{(i)} Z_i Z_{i+1} \right) \quad (1)$$

where X, Y, Z are the standard Pauli operators and $J_\alpha^{(i)}$ are bond-resolved coupling strengths. In the isotropic limit ($J_x = J_y = J_z$) the model possesses an exact $\text{SU}(2)$ symmetry; for $J_x = J_y \neq J_z$ (XXZ model) the $\text{U}(1)$ symmetry generated by $\sum Z_i$ is preserved, while for generic couplings (XYZ model) this symmetry is broken to $\mathbb{Z}_2 \times \mathbb{Z}_2$ [15]. In HVK, we fix the network context to six sites and treat every component $J_\alpha^{(i)}$ as a learnable real parameter, initialized from:

$$J_\alpha^{(i)} \sim \mathcal{N}(0, 0.01), \quad \alpha \in \{x, y, z\}. \quad (2)$$

Because the coupling strengths are optimized jointly with the circuit weights, the Hamiltonian itself adapts to image data, and the effective spin-spin interaction pattern is selected by the model dynamically.

2) *Energy Loss and Ground-State Objective:* Given a patch n and the Pauli observables measured from the variational circuit, the Heisenberg energy of the corresponding state is:

$$\mathcal{E}_n(\theta, J) = \sum_{i=0}^4 \sum_{\alpha \in \{x, y, z\}} J_\alpha^{(i)} \langle \sigma_\alpha^{(i)} \sigma_\alpha^{(i+1)} \rangle_n. \quad (3)$$

where α is the indices running from 1 to 3 denoting the pauli matrices.

The training energy loss is evaluated as the patch-averaged energy:

$$\mathcal{L}_{\text{energy}}(\theta, J) = \langle E \rangle = \frac{1}{N} \sum_{n=1}^N \mathcal{E}_n(\theta, J). \quad (4)$$

Minimizing $\mathcal{L}_{\text{energy}}$ with antiferromagnetic couplings drives the circuit toward the ground state of the learned Hamiltonian, steering the architecture toward spatially correlated quantum features.

3) *Positional Awareness of the Hamiltonian:* The Hamiltonian is rendered spatially aware through two coupled mechanisms. First, the per-qubit positional rotation θ_{position} depends on the patch-grid coordinate through a learned projection, so distinct spatial positions inject different inputs into the same Hamiltonian topology. Second, because the circuit output state $|\Psi_n\rangle$ depends jointly on the patch feature and the learned positional angle, the effective energy surface experienced by each patch is distinct even when underlying internal features align.

4) *Total Training Objective:* The complete objective function combines pixel-level fidelity with quantum-state regularization:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{rec}} + \lambda \mathcal{L}_{\text{energy}}, \quad (5)$$

evaluated via the Frobenius-norm reconstruction loss:

$$\mathcal{L}_{\text{rec}} = \frac{1}{N} \sum_{n=1}^N \|D(o_n, e_n) - P_n\|_F^2 \quad (6)$$

where D is the classical decoder, P_n is the ground-truth patch, and λ is the energy regularization coefficient.

5) *Symmetry Analysis:* The Heisenberg correlators define three observable sectors:

$$\Xi_{ZZ} = (\langle Z_0 Z_1 \rangle, \dots, \langle Z_4 Z_5 \rangle), \quad \Xi_{XX}, \quad \Xi_{YY} \quad (7)$$

embedded in $\mathbb{R}^{27}$ through the simultaneously measured observables. Two symmetry reductions act sequentially on this manifold:

- 1) **Positional reduction:** The per-qubit positional rotations belong to $\text{SU}(2)$ (acting as a double cover to the spatial $\text{SO}(3)$ Bloch rotations) and break permutation symmetry between qubit sites, reducing continuous symmetry down to the diagonal phase rotations matching the Cartan sector.
- 2) **Geometric reduction:** The 4×4 patch grid carries the dihedral group of the square,

$$D_4 = \langle r, s \mid r^4 = s^2 = (sr)^2 = e \rangle, \quad (8)$$

with eight elements. Under transformation matrices $g \in D_4$, observable vectors transform as $o_{g(i,j)} \approx \Phi_g o_{i,j}$, decomposing the latent space into irreducible representations:

$$\mathcal{V}_{\text{obs}} \approx m_A A_1 \oplus m_B A_2 \oplus m_C B_1 \oplus m_D B_2 \oplus m_E E. \quad (9)$$

The resulting symmetry stack handles the structural degradation systematically:

$$\text{SU}(2)\text{-variant isotropic sector} \xrightarrow{\text{positional encoding}} \text{Cartan sector} \xrightarrow{\text{grid geometry}} L \quad (10)$$

This framework preserves the structural symmetries of the patch grid, allowing the pipeline to distinguish orientation-sensitive boundaries.

III. HVK ARCHITECTURE DESIGN

The pipeline consists of an operational flow defined by:

Image $\rightarrow \{\text{Patches}\} \rightarrow \text{MPS} \rightarrow 46\text{-D Feature} \rightarrow 6\text{-D Vector} \rightarrow \text{Initial Latent} \rightarrow \text{Decoder} \rightarrow \text{Reconstruction}$

A. Tensor Network Optimization

1) *Patch Preparation*: An input image $\mathbf{I} \in \mathbb{R}^{H \times W}$ is normalized and resized to a canonical resolution of 256×256 pixels. It is partitioned into a regular grid of non-overlapping 64×64 patches:

$$\Pi = \{P_{ij} \in \mathbb{R}^{64 \times 64} : i, j \in \{0, 1, 2, 3\}\}, \quad (11)$$

yielding $N = 16$ patches. Each patch carries a normalized center coordinate $(r_i, c_j) = (i/4, j/4) \in \mathbb{R}^2$.

2) *MPS Feature Extraction*: Each patch P_{ij} is unrolled into a flat vector $\mathbf{v}_n$, ℓ_2 -normalized, and reshaped into a 12-site spin tensor product state:

$$|\psi_n\rangle = \frac{\mathbf{v}_n}{\|\mathbf{v}_n\|_2} \in (\mathbb{C}^2)^{\otimes 12}, \quad (12)$$

satisfying $2^{12} = 4096$ variables. The state is decomposed into a Matrix Product State via sequential Singular Value Decomposition (SVD) with a strict rank truncation set to bond dimension $\chi = 4$:

$$|\psi_n\rangle \approx \sum_{\sigma_1, \dots, \sigma_{12}} A_{\sigma_1} A_{\sigma_2} \dots A_{\sigma_{12}} |\sigma_1 \dots \sigma_{12}\rangle, \quad (13)$$

where $A_{\sigma_i} \in \mathbb{C}^{\chi_{i-1} \times \chi_i}$. This provides a genuine quantum-aligned feature envelope protecting local correlations.

3) *Observable Feature Extraction*: Three distinct measurement families are extracted from the underlying MPS to constitute a 46-dimensional classical feature vector $\mathbf{f}_n$:

- **Local Pauli Expectations**: $\langle Z_i \rangle$ and $\langle X_i \rangle$ for all 12 sites, contributing 24 features.
- **Nearest-Neighbor Correlations**: $\langle Z_i Z_{i+1} \rangle$ across 11 spatial bonds, contributing 11 features.
- **Bipartite Entanglement Entropy**: Calculated across all 11 split configurations via singular values $\lambda_k^{(i)}$:

$$p_k^{(i)} = \frac{|\lambda_k^{(i)}|^2}{\sum_k |\lambda_k^{(i)}|^2}, \quad S_i = - \sum_k p_k^{(i)} \log(p_k^{(i)} + \varepsilon), \quad (14)$$

with $\varepsilon = 10^{-8}$, yielding 11 features.

The final combined configuration produces $\mathbf{f}_n \in \mathbb{R}^{46}$.

4) *Two-Dimensional Fourier Positional Encoding*: Spatial information is injected via a sinusoidal map to break coordinate invariance [17]. For coordinate pair (r_i, c_j) , the vector $e \in \mathbb{R}^8$ is defined by:

$$\begin{aligned} e_{4k} &= \sin(\omega_k \pi r_i), & e_{4k+1} &= \cos(\omega_k \pi r_i), \\ e_{4k+2} &= \sin(\omega_k \pi c_j), & e_{4k+3} &= \cos(\omega_k \pi c_j), \end{aligned} \quad (15)$$

for $k \in \{0, 1\}$ with base frequencies ω_k . This is projected to rotational constraints via a learned map $W_{\text{pos}} \in \mathbb{R}^{6 \times 8}$:

$$\theta_{\text{position}} = W_{\text{pos}} \cdot e \in \mathbb{R}^6. \quad (16)$$

5) *Variational Quantum Circuit*: The raw 46-D feature matrix $\mathbf{f}_n$ is linearly compressed to a 6-D vector $x_n \in \mathbb{R}^6$ prior to circuit injection. For each patch, a 6-qubit register initialized to $|0\rangle^{\otimes 6}$ undergoes the following operations:

- 1) **Angle Embedding**: Feature inputs are mapped via Y -axis rotations:

$$U_{\text{enc}}(x_n) = \bigotimes_{i=0}^5 R_Y(x_n^{(i)}), \quad R_\theta = e^{-i\theta Y/2}. \quad (17)$$

- 2) **Positional Modulation**: Spatial adjustments are handled via:

$$U_{\text{pos}}(\theta_{\text{pos}}, n) = \bigotimes_{i=0}^5 R_Z(\theta_{\text{position}}^{(i)}). \quad (18)$$

- 3) **Strongly Entangling Layers**: Parameterized operations use $L = 2$ layers containing single-qubit rotations and CNOT rings:

$$U_{\text{var}}(\theta) = \prod_{l=1}^L \left[\text{CNOT-Ring} \cdot \bigotimes_{i=0}^5 \text{Rot}(\phi_i^l, \theta_i^l, \omega_i^l) \right], \quad (19)$$

initialized safely to avoid barren plateaus [1]. The state is output as:

$$|\Psi_n\rangle = U_{\text{var}}(\theta) \cdot U_{\text{pos}}(\theta_{\text{pos}}, n) \cdot U_{\text{enc}}(x_n) |0\rangle^{\otimes 6}. \quad (20)$$

6) *Observable Latent Space*: A total of 27 real expectations are evaluated per execution pass:

$$\Phi_n = (\langle X_i \rangle, \langle Z_i \rangle, \langle Z_i Z_{i+1} \rangle, \langle X_i X_{i+1} \rangle, \langle Y_i Y_{i+1} \rangle) \in [-1, 1]^{27}, \quad (21)$$

representing a structured projection $\Phi : \mathbb{CP}^{63} \rightarrow \mathbb{R}^{27}$.

7) *Classical Decoding and Patch Reconstruction*: The extracted matrix is standardized to standard normal form per column and passed to a classical multi-layer perceptron (MLP) decoder alongside the spatial encodings:

$$\tilde{P}_n = D(\text{Standardize}(\Phi_n), e_n). \quad (22)$$

The architecture utilizes a mirrored transpose structure with expanding channel maps ($4 \rightarrow 8 \rightarrow 16 \rightarrow 3$) coupled with cubic upsampling steps to output the final pixel arrays [20].

IV. METHODOLOGY

A. Datasets

1) *Monalisa Single-Image Benchmark*: The baseline canonical single-image analysis uses a 256×256 grayscale layout of *monalisa.jpg*, divided into 16 evaluation regions of size 64×64 pixels [16].

2) *CIFAR-10 Subset*: To capture structural generalization capabilities, a target subset of 10 images from the CIFAR-10 dataset is resized to match the 256×256 input dimensions. Performance tracking checks aggregate values across metrics repositories [18].

B. Benchmarking Implementations

We execute validation passes using four distinct comparison configurations:

- **1D HVK:** A control setup utilizing a linear chain of 6 qubits linked via 5 nearest-neighbor bonds to generate the 27-dimensional observable vector.
- **2D HVK:** An advanced geometric layout arranging the 6 qubits on a 2×3 lattice to match patch boundary configurations natively, producing a 19-dimensional latent signature.
- **Convolutional GAN Autoencoder:** A standard generative adversarial network processing matching patch blocks [18].
- **Parameterized Hamiltonian Learning (PHL):** A task-mismatched segmentation reference used to establish positioning parameters rather than structural metrics [22].

C. Training Protocol

Optimization passes utilize the Adam optimizer with learning rates set to $\eta = 0.003$ for the 1D variant and $\eta = 0.004$ for the 2D configuration [18]. Parameter targets are jointly updated over 200 epochs under a rigid loss factor ($\lambda = 0.01$). Initialization parameters scale according to established safe limits [1], [16], [17].

D. Evaluation Metrics

1) *Reconstruction Precision:* The system gauges reconstruction performance using Mean Squared Error (MSE) and Peak Signal-to-Noise Ratio (PSNR):

$$\text{MSE}_A(t) = \frac{1}{N \cdot 64^2} \sum_{n=1}^N \|\tilde{P}_{n,t}^{(A)} - P_n\|_F^2, \quad (23)$$

$$\text{PSNR}_A(t) = 20 \log_{10} \left(\frac{1}{\sqrt{\text{MSE}_A(t)}} \right). \quad (24)$$

2) *Hamiltonian Energy Tracking:* The state optimization depth is directly tracked via the underlying Heisenberg expectation tracking:

$$\langle E \rangle_t = \frac{1}{N} \sum_{n=1}^N \sum_i \left(J_z^{(i)} \langle Z_i Z_{i+1} \rangle_{n,t} + J_x^{(i)} \langle X_i X_{i+1} \rangle_{n,t} + J_y^{(i)} \langle Y_i Y_{i+1} \rangle_{n,t} \right). \quad (25)$$

3) *Order Parameters and Phase Diagnostics:* Longitudinal order parameters track system magnetization states alongside lattice correlation configurations:

$$C_A(t) = \frac{1}{15 \cdot N} \sum_{n=1}^N \sum_{i=0}^4 (\langle Z_i Z_{i+1} \rangle_{n,t} + \langle X_i X_{i+1} \rangle_{n,t} + \langle Y_i Y_{i+1} \rangle_{n,t}). \quad (26)$$

Critical structural transition markers are flagged using windowed variance deviations:

$$|C_A(t_c) - \text{median}(C_A)| > 2 \cdot \text{std}(C_A). \quad (27)$$

TABLE I
CIFAR-10 SUBSET EVALUATION (10-IMAGE AVERAGE)

Metric	1D HVK	GAN Baseline	Direction
Final MSE	6.7974×10^{-5}	4.1511×10^{-4}	Lower is better
Best MSE	6.1102×10^{-5}	3.8944×10^{-4}	Lower is better
PSNR (dB)	42.55	34.03	Higher is better
SSIM	—	0.9944	Higher is better
Final Energy	-0.1844	—	Lower is better
Critical Epoch	26.1	—	—

V. EXPERIMENTAL RESULTS

A. CIFAR-10 Subset Results

Table I compiles the structural reconstruction performance metrics evaluated across the CIFAR-10 data configurations.

The 1D HVK tracks an 83.6% reduction in absolute mean squared error compared directly to the GAN alternative, maintaining a clean +8.52 dB performance advantage.

B. Monalisa Single-Image Results

Reconstruction comparisons for the structural *Monalisa* validation target are summarized in Table II.

TABLE II
MONALISA RECONSTRUCTION PERFORMANCE ACROSS METHODS

Model	Final MSE	Best MSE	PSNR (dB)	SSIM
1D HVK	3.3754×10^{-4}	3.0112×10^{-4}	34.72	—
2D HVK	8.5114×10^{-5}	8.4975×10^{-5}	40.70	—
GAN Baseline	1.0185×10^{-3}	1.1124×10^{-3}	29.92	0.9873
PHL [22]	—	—	—	—

The PHL metric is excluded from direct pixel tracking because it is natively optimized for segmentation tasks, logging a baseline Dice coefficient of 0.9706 and an Intersection-over-Union (IoU) of 0.9428.

C. 1D vs. 2D Topological Ablation Study

Table III highlights the operational adjustments caused by restructuring the lattice configuration from a 1D chain to a 2D adjacent grid.

TABLE III
LATTICE TOPOLOGY PERFORMANCE ADJUSTMENTS

Metric	1D HVK	2D HVK
Critical Epoch	11	2
Final Energy	-0.1154	-0.2411
Order Parameter $\langle M_z \rangle$	0.0411	0.0094
Maximum Susceptibility $\chi_{\max}$	1.15×10^{-4}	6.24×10^{-4}
Final PSNR (dB)	34.72	40.70

D. Quantum Contribution Ablation Study

The central empirical question is not whether HVK can reconstruct a single target image, but which part of the hybrid pipeline is responsible for that reconstruction. We therefore ran a controlled ablation suite on the *Monalisa* benchmark

using the same image size, patching rule, seed, and metric definitions whenever possible. Tables IV–VI summarize the completed experiments. Lower MSE and higher PSNR/SSIM are better.

The evaluation controls show that the decoder is not reconstructing solely from patch coordinates. Shuffling observable vectors across spatial positions drops PSNR from 32.24 dB to 19.70 dB, and replacing observables with zeros while retaining real positional encodings drops PSNR to 14.85 dB. Thus, the Pauli-observable channel carries patch-specific information. At the same time, the freeze and replacement controls impose a sharp qualification: freezing the VQC weights and Hamiltonian couplings does not harm reconstruction, and purely classical replacements match or exceed the trained VQC baseline. These results support the claim that the *observable-latent bottleneck* is useful, but they do not support a claim that trained quantum parameters, entangling gates, or quantum advantage are responsible for the measured reconstruction quality.

The Hamiltonian controls are the most important corrective result in the study. The legacy objective adds a signed mean energy term to the reconstruction loss. Because the learned energy can become negative, the optimizer can lower total loss by driving energy downward without improving reconstruction. Empirically, removing this term improves PSNR from 32.24 dB to 33.30 dB. We also tested a contrastive Hamiltonian objective that assigns lower energy to correct feature–position pairs than to shuffled pairs; this improves over the legacy baseline but remains below its own no-energy-loss counterpart. Consequently, the Hamiltonian should be interpreted as an energy diagnostic and correlation probe in the present system, not as the primary driver of reconstruction accuracy.

The step-count sweep confirms that the model is convergence-limited at early iteration counts and approaches a plateau near 500 steps. Within the tested circuit widths, $q = 8$ gives the strongest qubit-count result. The MPS bond-dimension sweep is more nuanced: $\chi = 1$ and $\chi = 2$ outperform the default $\chi = 4$ under the same training budget, while $\chi = 8$ partially recovers. This suggests that reconstruction quality is not controlled only by MPS compression fidelity; downstream projection, circuit width, and decoder training dynamics interact with the feature basis.

The generalization controls clarify scope. A model optimized on *Monalisa* does not reconstruct a structurally different image without retraining, producing only 7.78 dB PSNR. When training is extended to include the second image, PSNR rises to 28.31 dB with SSIM 0.9862. HVK should therefore be presented as a per-image or small-set reconstruction framework at this stage; broad dataset-level generalization remains a future validation target.

VI. DISCUSSION

A. Reconstruction Error Analysis

HVK’s 83.6% error reduction over the tested GAN configuration demonstrates that adversarial objectives can introduce structural instabilities during high-fidelity generation. GAN optimization exhibits a non-monotonic MSE trace ($\text{MSE}_{\text{best}} >$

$\text{MSE}_{\text{final}}$), whereas HVK uses direct pixel losses and an observable bottleneck whose patchwise alignment can be explicitly ablated. The completed ablations show that this observable channel is meaningful, but also that Hamiltonian energy regularization is not the current source of the reconstruction gain.

B. Architectural Differentiation Matrix

Table VIII contrasts the structural properties of HVK against current reference frameworks.

C. The Topology-Alignment Principle

The reported +5.98 dB performance advantage of the 2D HVK configuration over its 1D counterpart supports a *topology-alignment principle*: aligning the latent interaction graph with image-patch adjacency can improve reconstruction. The ablation study tempers this claim by showing that circuit topology alone is not sufficient evidence of quantum advantage; topology alignment should be treated as an architectural inductive bias whose contribution must be separated from decoder capacity.

D. Quantitative Channel Economy

The application of 6 qubits maps precisely to 27 or 19 target measurement dimensions, tracking the specific number of local and nearest-neighbor interactions. HVK maintains compact scaling rules while outpacing over-parameterized convolutional models.

E. Interpretive Limits

We explicitly acknowledge the following bounds on our empirical claims:

- The CIFAR-10 evaluations are currently bounded to a 10-image validation subset.
- 2D HVK performance data is documented solely for the *Monalisa* test case.
- The single-image *Monalisa* experiments are per-image optimization studies and should not be interpreted as dataset-level generalization.
- Current ablations show that observable latents are load-bearing, but they do not establish that trained VQC parameters, entangling gates, or the Heisenberg energy term are the unique source of reconstruction quality.
- CIFAR-10 benchmarking indicates that HVK is competitive against several baselines, including the tested GAN, PHL, and classical autoencoder configurations, but high-capacity MLP and CNN baselines remain stronger in pure reconstruction accuracy.

VII. CONCLUSION

We introduced HVK, a tensor-network and quantum-observable framework for spatially aware image reconstruction. HVK combines an MPS encoder, positional inputs, a parameterized variational circuit, Pauli-observable latent vectors, and Hamiltonian energy diagnostics.

TABLE IV

COMPONENT-LEVEL ABLATIONS ON THE *Monalisa* BENCHMARK. THE OBSERVABLE CHANNEL IS LOAD-BEARING, BUT TRAINED VQC PARAMETERS ARE NOT ISOLATED AS THE DOMINANT SOURCE OF PERFORMANCE.

Experiment	MSE	PSNR (dB)	SSIM	Interpretation
Shared VQC baseline	5.9771×10^{-4}	32.24	0.9919	Reference standard HVK run
Shuffle observables at eval	1.0727×10^{-2}	19.70	0.8423	Observable–position alignment matters
Zero observables, real positions	3.2754×10^{-2}	14.85	0.6431	Position alone is insufficient
Freeze classical, train quantum only	8.0639×10^{-2}	10.93	0.0214	Frozen decoder/projections collapse
Freeze quantum, train classical only	4.8819×10^{-4}	33.11	0.9934	Fixed quantum features remain usable
Random VQC output	2.5230×10^{-3}	25.98	0.9656	Random latents are substantially worse
No entanglement	5.4720×10^{-4}	32.62	0.9926	Entangling gates are not decisive here
Classical tanh replacement	4.5146×10^{-4}	33.45	0.9939	Classical map matches or exceeds VQC
Parameter-matched classical map	4.7968×10^{-4}	33.19	0.9935	Small classical control also exceeds baseline

TABLE V

HAMILTONIAN AND OBSERVABLE-SECTOR CONTROLS. THE CURRENT HAMILTONIAN ENERGY OBJECTIVE IS DIAGNOSTICALLY USEFUL BUT DOES NOT IMPROVE RECONSTRUCTION ACCURACY.

Experiment	MSE	PSNR (dB)	SSIM	Interpretation
Legacy baseline, signed energy	5.9771×10^{-4}	32.24	0.9919	Standard energy-regularized run
No energy loss	4.6772×10^{-4}	33.30	0.9937	Removing energy improves PSNR
No observable noise	5.3518×10^{-4}	32.72	0.9928	Noise is not essential
Contrastive Hamiltonian core	5.1975×10^{-4}	32.84	0.9930	Stronger Hamiltonian objective helps modestly vs. baseline
Contrastive, no energy loss	4.6498×10^{-4}	33.33	0.9937	Energy removal still performs best
ZZ-only observables	5.1503×10^{-4}	32.88	0.9930	Nearest-neighbor ZZ correlations are largely sufficient

TABLE VI

CAPACITY AND FEATURE-BOTTLENECK SWEEPS. INCREASING QUBIT COUNT IMPROVES PERFORMANCE WITHIN THE TESTED RANGE, WHILE THE MPS BOND-DIMENSION SWEEP DOES NOT SHOW MONOTONIC GAINS.

Experiment	Setting	MSE	PSNR (dB)	SSIM
Step sweep	30 steps	1.4777×10^{-2}	18.30	0.7467
Step sweep	60 steps	4.8797×10^{-3}	23.12	0.9306
Step sweep	120 steps	1.3336×10^{-3}	28.75	0.9818
Step sweep	240 steps	5.5542×10^{-4}	32.55	0.9927
Step sweep	500 steps	4.5941×10^{-4}	33.38	0.9938
Qubit count	$q = 4$	5.3633×10^{-4}	32.71	0.9928
Qubit count	$q = 6$	5.9771×10^{-4}	32.24	0.9919
Qubit count	$q = 8$	4.9322×10^{-4}	33.07	0.9933
MPS bond dim.	$\chi = 1$	5.0746×10^{-4}	32.95	0.9932
MPS bond dim.	$\chi = 2$	5.3061×10^{-4}	32.75	0.9929
MPS bond dim.	$\chi = 4$	5.9771×10^{-4}	32.24	0.9919
MPS bond dim.	$\chi = 8$	5.6124×10^{-4}	32.51	0.9925
No MPS, patch statistics	—	5.3874×10^{-4}	32.69	0.9927

TABLE VII

GENERALIZATION CONTROLS. SINGLE-IMAGE OPTIMIZATION DOES NOT TRANSFER ZERO-SHOT, WHILE MULTI-IMAGE TRAINING GIVES A USEFUL BUT WEAKER RECONSTRUCTION ON THE SECOND IMAGE.

Experiment	MSE	PSNR (dB)	SSIM
Second image, zero-shot	1.6657×10^{-1}	7.78	0.0196
Second image, multi-image training	1.4740×10^{-3}	28.31	0.9862

On CIFAR-10, HVK achieves a PSNR of 42.55 dB in the reported 10-image subset, outperforming the tested GAN baseline by 8.52 dB and reducing mean squared error by 83.6%. On the *Monalisa* target, the 2D HVK setup delivers a PSNR of 40.70 dB, exceeding the 1D control by 5.98 dB and the GAN by 10.78 dB. These results indicate that the observable-latent reconstruction pipeline is competitive with several non-convolutional reconstruction baselines.

However, the quantum-contribution ablation study imposes

an important qualification. The observable channel is load-bearing, but the current experiments do not demonstrate that the trained VQC, entangling gates, or Hamiltonian energy loss are the decisive source of performance. Freeze-quantum, no-entanglement, no-energy-loss, and parameter-matched classical replacement controls all reconstruct at or above the shared VQC baseline. HVK should therefore be framed as a hybrid observable-latent reconstruction architecture and experimental platform, rather than as evidence of quantum advantage in its current form.

Future efforts will focus on full dataset-level train/test splits, multiple random seeds, stricter parameter-matched classical controls, reduced-capacity decoders, and tasks where Hamiltonian structure, entanglement, or observable dynamics can be isolated as a genuine source of generalization or regularization.

REFERENCES

- [1] M. Cerezo et al., “Variational quantum algorithms,” *Nature Reviews Physics*, vol. 3, no. 9, pp. 625–644, 2021.

TABLE VIII
ARCHITECTURAL DIFFERENTIATION MATRIX

Axis	Classical AE [20]	GAN AE [18]	Standard QAE [13]	HVK (Ours)
Latent Space	Continuous $\mathbb{R}^d$	Continuous $\mathbb{R}^d$	Quantum Register	Pauli Correlator Vector
Spatial Bias	Implicit Convolutions	Implicit Convolutions	None	2D Fourier + D_4 Symmetries
Regularizer	Bottleneck Size	Discriminator Loss	Swap-Test Fidelity	Optional Hamiltonian energy diagnostic
Group Alignment	None	None	Reference Swaps	SO(3) Cartan + D_4 Grid
Target Task	Reconstruction	Reconstruction	Compression	Spatially Aware Reconstruction

- [2] J. Tilly et al., “The Variational Quantum Eigensolver: A review of methods and best practices,” *Physics Reports*, vol. 986, pp. 1–128, 2022.
- [3] B. Bauer, S. Bravyi, M. Motta, and G. K.-L. Chan, “Quantum Algorithms for Quantum Chemistry and Quantum Materials Science,” *Chemical Reviews*, vol. 120, no. 22, pp. 12685–12717, 2020.
- [4] S. Chakraborty, S. B. Mandal, and S. H. Shaikh, “Quantum image processing: challenges and future research issues,” *International Journal of Information Technology*, vol. 14, no. 1, pp. 475–489, 2018.
- [5] R. Zhou, W. Hu, P. Fan, and H. Ian, “Quantum realization of the bilinear interpolation method for NEQR,” *Scientific Reports*, vol. 7, no. 1, p. 2511, 2017.
- [6] X. Yao et al., “Quantum Image Processing and Its Application to Edge Detection: Theory and Experiment,” *Physical Review X*, vol. 7, no. 3, p. 031041, 2017.
- [7] J. I. Cirac, D. Pérez-García, N. Schuch, and F. Verstraete, “Matrix product states and projected entangled pair states: Concepts, symmetries, theorems,” *Reviews of Modern Physics*, vol. 93, no. 4, p. 045003, 2021.
- [8] J. I. Latorre, “Image compression and entanglement,” arXiv:quant-ph/0510031, 2005.
- [9] Z. Han, J. Wang, H. Fan, L. Wang, and P. Zhang, “Unsupervised Generative Modeling Using Matrix Product States,” *Physical Review X*, vol. 8, no. 3, p. 031012, 2018.
- [10] S. M. Fei, “Compressed Sensing Based on Tensor Network Machine Learning,” *Physical Science & Biophysics Journal*, vol. 5, no. 1, p. 16000174, 2021.
- [11] D. Guala et al., “Practical overview of image classification with tensor-network quantum circuits,” *Scientific Reports*, vol. 13, no. 1, p. 4427, 2023.
- [12] A. Haga, “Quantum annealing-based computed tomography using variational approach for a real-number image reconstruction,” *Physics in Medicine and Biology*, vol. 69, no. 4, p. 02155, 2024.
- [13] X. Zhai, T. Xiao, J. Huang, J. Fan, and G. Zeng, “Quantum neural compressive sensing for ghost imaging,” *Physical Review Applied*, vol. 23, no. 1, p. 014018, 2025.
- [14] M. A. Nau, A. H. Vija, W. Gohn, M. P. Reymann, and A. Maier, “Exploring the Limitations of Hybrid Adiabatic Quantum Computing for Emission Tomography Reconstruction,” *Journal of Imaging*, vol. 9, no. 10, p. 221, 2023.
- [15] B. Craps, M. D. Clerck, O. Evnin, and P. Hacker, “Integrability and complexity in quantum spin chains,” *SciPost Physics*, vol. 16, no. 2, p. 041, 2024.
- [16] “Mathematical Framework for HVK,” internal technical report, June 2026.
- [17] R. Xu, X. Wang, K. Chen, B. Zhou, and C. C. Loy, “Positional Encoding as Spatial Inductive Bias in GANs,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021, pp. 13564–13573.
- [18] “Methodology and Experimental Analysis of the Hamiltonian Vision Kernel,” internal laboratory reference, 2026.
- [19] J. Y. Araz and M. Spannowsky, “Classical versus quantum: Comparing tensor-network-based quantum circuits on Large Hadron Collider data,” *Physical Review A*, vol. 106, no. 6, p. 062423, 2022.
- [20] D. Slabbert and F. Petruccione, “Hybrid Quantum-Classical Feature Extraction approach for Image Classification using Autoencoders and Quantum SVMs,” arXiv:2410.18814, 2024.
- [21] Z. Liu, P.-X. Shen, W. Li, L.-M. Duan, and D.-L. Deng, “Quantum capsule networks,” *Quantum Science and Technology*, vol. 8, no. 1, p. 015016, 2022.
- [22] J. Shi, W. Wang, X. Lou, S. Zhang, and X. Li, “Parameterized Hamiltonian Learning With Quantum Circuit,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 45, no. 5, pp. 6086–6095, 2022.