

Hamiltonian Vision Kernel: Symmetry-Pooled Quantum Correlator Features with Hardware Validation

Sparsho Chakraborty* and Siddhartha Patra†

Abstract—We introduce the Hamiltonian Vision Kernel (HVK), a new hybrid quantum–classical architecture for image reconstruction. HVK assembles matrix-product-state patch features, a variational quantum circuit exposing a Pauli-correlator latent space, two-dimensional Fourier positional encoding, a learnable Hamiltonian energy diagnostic, and a classical decoder into a single pipeline; HVK1D and HVK2D instantiate chain and grid interaction graphs, and observable pooling gives the HVK2D map D_4 equivariance by construction. This specific combination is new, not any one ingredient in isolation (Section II). Across six real image datasets, per-image-trained HVK2D models reconstruct with high fidelity (25.8–41.6 dB PSNR). On a resource-matched, multi-seed held-out CIFAR-10 comparison, the tested HVK2D map does not exceed simple local-observable or raw-linear classical controls (18.12 vs. 18.80 dB); a pre-declared Two One-Sided Tests (TOST) equivalence test at a ± 1 dB margin confirms this gap is small enough to call the two *competitive* rather than merely indistinguishable, while HVK2D still beats a random-VQC and a strict-classical-random-feature control by a wide margin. HVK also executes on real IBM Quantum hardware: five fitted checkpoints are replayed directly from hardware-measured observables, with no retraining, reaching 25.90–31.52 dB PSNR. If HVK ties resource-matched classical baselines on ordinary reconstruction, its case rests elsewhere: an entangling observable channel that is representationally necessary on a leakage-audited task requiring distant-patch products (mean $R^2 = 0.974$ vs. $R^2 \leq 0.02$ for non-entangling controls), an interpretable Hamiltonian energy diagnostic, and native execution on quantum hardware – capabilities a purely classical reconstruction pipeline does not provide. A companion supplementary study additionally reports an exploratory training-dynamics change-point diagnostic (a magnetization-style statistic and an energy-to-entanglement ratio, varied along qubit-count and MPS-bond-dimension axes and partly replayed on real hardware); we report it as a scoped diagnostic rather than a headline result because its own negative control fires identically whether or not the Hamiltonian term is present. We make no claim of quantum advantage: HVK is presented as a new, working architecture that is competitive with resource-matched classical baselines, validated on real hardware, and equipped with capabilities a classical reconstruction pipeline does not natively expose. A companion supplementary study provides the resource-matched held-out controls, leakage audits, statistical tests, and complete hardware methodology behind these claims.

Index Terms—Quantum Image Processing, Tensor Networks, Variational Quantum Algorithms, Pauli Observables, Hamiltonian Diagnostics, Group-Equivariant Representations, Hybrid Quantum Hardware Validation.

I. INTRODUCTION

Hybrid quantum–classical models for vision tasks are increasingly proposed as a route to representations with favorable inductive bias, exploiting entanglement, physically motivated regularizers, or tensor-network structure that classical convolutional pipelines do not natively expose [1], [2], [20]. The variational quantum eigensolver and its descendants exploit the local decomposability of physically relevant Hamiltonians to construct shallow, problem-aware circuits [1]–[3], and quantum image processing provides encodings that place pixel information directly into qubit registers [4]–[6]. Tensor-network methods, in parallel, have proven effective as compressed representations of structured classical data, including images [7]–[11], [22], and more broadly other structured data such as particle-physics events [16]. Related hybrid quantum architectures target adjacent imaging tasks, including quantum annealing-based tomographic reconstruction [12], adiabatic quantum computing for emission tomography [14], quantum neural compressive sensing for ghost imaging [13], hybrid autoencoder-plus-quantum-SVM feature extraction for image classification [17], and quantum capsule networks [18]. HVK focuses on the specific combination of a Pauli-correlator latent space, enforced group-equivariant pooling, and a physically motivated energy diagnostic.

We introduce the Hamiltonian Vision Kernel (HVK1D/HVK2D), which combines these ingredients into a single architecture with three properties studied jointly here: a *Pauli-correlator latent space* (measured VQC observables, not a raw quantum state, form the representation handed to a classical decoder, in the spirit of quantum-enhanced feature maps [21]); an *exact, numerically-verified D_4 -equivariant observable-pooling construction*, group-averaging the observable grid over the eight symmetries of the square rather than merely motivating symmetry through architecture choices, in the spirit of group-equivariant classical vision [23] and symmetry-embedding constructions for variational circuits [24]; and a *learnable Hamiltonian energy diagnostic* (full nearest-neighbor $XX/YY/ZZ$ sectors in HVK1D and grid-edge ZZ terms in HVK2D). Table VIII contrasts these properties against classical, adversarial, and standard quantum autoencoder baselines.

a) Novelty relative to the nearest hybrid-vision architectures.: HVK’s closest architectural relatives are three hybrid quantum-vision patterns already established individually in the literature, and situating HVK against each clarifies what is

*School of Basic Sciences, Department of Physics, Indian Institute of Technology Bhubaneswar, Odisha, India.

†Assistant Professor, Centre for Quantum Engineering, Research and Education (CQUERE), TCG CREST, Kolkata, India.

new here versus what is inherited. *Register-based quantum autoencoders* [28]–[30] encode an image’s pixel amplitudes directly into a qubit register and train a trash-qubit fidelity objective (or attach a classical decoder to that register’s measured amplitudes); HVK never amplitude-encodes pixels into qubits at all — it exposes only classically pre-compressed MPS patch features as circuit input, and hands the decoder measured Pauli correlators rather than register amplitudes or a fidelity score. *Quantum-kernel vision models* [21], [31] embed an input into a quantum feature space and read out a scalar kernel (state-overlap) value consumed by a classical kernel method; HVK instead reads out an explicit vector of local and pairwise Pauli expectation values as a latent representation for a trained decoder, and adds enforced group-equivariant pooling and a Hamiltonian regularizer that a kernel construction has no analogue for. *Tensor-network (TN) circuit classifiers* [9], [11], [22] use MPS structure either as a standalone classical model or as a template for a shallow quantum circuit trained for classification; HVK uses MPS compression only as a classical *preprocessing* step that sets a separately trained VQC’s rotation angles, and targets per-image reconstruction rather than classification. None of these three families combines classical MPS-compressed input, a measured Pauli-correlator (not kernel- or register-state) latent space, enforced D_4 -equivariant pooling, and a learnable Hamiltonian diagnostic in one pipeline; HVK’s contribution is this specific *assembly* and the properties it jointly enables (Table VIII), not any single ingredient in isolation — each of which individually recurs throughout this literature (Section V-C).

b) Contributions.: **1) Architecture:** HVK1D/HVK2D combines MPS patch features, a six-qubit VQC, a 19–27-dimensional Pauli-correlator vector, 2D Fourier positional encoding, a graph-Hamiltonian diagnostic, and a classical decoder.

2) Reconstruction scope: Identically configured per-image fitting experiments cover six real image datasets (Section IV-A).

3) Competitive with resource-matched classical baselines: A pre-declared TOST equivalence test on held-out CIFAR-10 shows HVK2D is statistically equivalent, within a ± 1 dB margin, to six of eight resource-matched classical/ablated controls, while still beating the two weakest controls by a wide margin (Section IV-C).

4) Hardware pilot: Five fitted checkpoints are replayed on IBM hardware without retraining and decoded from hardware-measured observables (Section IV-D), strengthened by a noise/shot-budget sweep and repeated execution across backends (Sections IV-E–IV-F).

5) Equivariant observable pooling: An exact, numerically-verified D_4 -equivariant observable-pooling construction makes the HVK2D observable map equivariant to the eight symmetries of the square, to floating-point precision (Section IV-G). The guarantee follows from Reynolds group averaging over the D_4 orbit, not from anything quantum: the identical construction applied to a purely classical feature map yields the same exactness (Supplementary Material). We therefore report it as a property the assembled pipeline provides, not as a quantum result.

6) Restricted nonlocal result: In a leakage-audited fixed-

readout protocol, entangling pair observables are necessary when the target depends explicitly on distant-patch products (Section IV-H).

The companion supplementary document additionally reports an exploratory training-dynamics change-point diagnostic (a magnetization-style statistic and an energy-to-entanglement ratio, tracked across qubit-count and MPS-bond-dimension axes and partly replayed on real hardware). We do not list it as a headline contribution: its own negative control shows the detector fires identically whether or not the Hamiltonian term is present, so it is reported as a scoped, self-critiqued diagnostic rather than evidence of a physical transition. A companion document, *Supplementary Material for the Hamiltonian Vision Kernel: Resource-Matched Ablations, Validation, and Reproducibility*, reports resource-matched held-out comparisons of HVK and classical/ablated feature maps, together with full hardware-pilot methodology (circuit construction, validation, job identifiers, quota accounting) and explicit provenance audits of excluded exploratory results. We keep that analysis in a companion document so the architecture-level contributions above and the generalization question do not obscure each other; Section IV-C summarizes its central finding and how it relates to the results here.

II. HAMILTONIAN VISION KERNEL: ARCHITECTURE

HVK treats each spatial patch as a finite spin state, compresses it as a matrix product state (MPS), measures a fixed set of Pauli observables from a variational quantum circuit (VQC), and decodes those observables back into pixels with a classical multi-layer perceptron (MLP). A learnable graph-Hamiltonian energy term regularizes training; unlike parameterized Hamiltonian learning approaches that treat the Hamiltonian as the primary learned object [19], HVK’s energy term is a diagnostic regularizer alongside a separate reconstruction objective. The shared pipeline is:

$$\begin{aligned}
 \text{Image} &\rightarrow \{\text{Patches}\} \rightarrow \text{MPS} \\
 &\rightarrow 46\text{-D}/22\text{-D Feature} \rightarrow 6\text{-D Vector} \\
 &\rightarrow 27\text{-D}/19\text{-D Latent} \rightarrow \text{Decoder} \\
 &\rightarrow \text{Reconstruction.}
 \end{aligned} \tag{1}$$

Three of these design choices warrant explicit motivation rather than being left implicit in the pipeline above. *Why MPS features:* treated as an amplitude-encoded quantum state, a patch of n sites has a coefficient tensor with 2^n entries; matrix-product-state factorization instead represents it as a chain of small tensors bounded by a fixed bond dimension χ , a compact, singular-value-controlled compression whose entanglement-entropy profile is already known to track image compressibility [8], [22], and whose local expectations and correlations are cheap classical features rather than requiring a full state-preparation circuit. *Why a Pauli-correlator latent space rather than a raw quantum state:* handing the classical decoder a fixed vector of measured Pauli expectation values, instead of full state tomography, keeps the quantum-to-classical interface at $O(\text{qubits})$ rather than exponential readout cost and follows the same measured-observable-as-feature pattern used by quantum-enhanced feature maps generally

[21]; it is also what makes the interpretability diagnostics of Section V-A and the training-dynamics diagnostics reported in the companion Supplementary Material possible, since a full state vector would not expose a comparably compact, physically meaningful summary. *Why a grid rather than chain interaction topology for HVK2D*: aligning the six-qubit interaction graph with the 2×3 patch-neighbor adjacency, rather than an arbitrary linear chain, is what makes the D_4 group-pooling construction of Section IV-G well-defined in the first place — the pooling operator averages over the spatial symmetries of the square, which requires a latent structure that itself respects the patch grid’s two-dimensional adjacency. HVK1D’s chain topology is retained as the ungridded control this comparison needs (Section V-D), not because a chain is separately motivated for 2D image data.

Two executed preprocessing configurations share this pipeline. For the 256×256 *Monalisa*/HVK1D benchmark, sixteen non-overlapping 64×64 patches are ℓ_2 -normalized, reshaped as 12-site spin states, and sequentially SVD-compressed to MPS bond dimension $\chi = 4$. Their 24 local X/Z expectations, 11 nearest-neighbor ZZ correlations, and 11 bipartite entropies form a 46-D feature vector. For the native 32×32 HVK2D multi-dataset experiments, sixteen 8×8 patches are represented by six-site MPSs, giving 12 local expectations, five nearest-neighbor ZZ terms, and five entropies (22 dimensions). A learned linear layer compresses either MPS feature vector to six circuit angles. HVK1D then measures 27 observables (six local Z , six local X , and five each of ZZ , XX , and YY); HVK2D measures 19 (six local Z , six local X , and seven grid-edge ZZ terms). The respective Fourier positional vectors have dimension eight and four, following the established use of Fourier positional encoding as a spatial inductive bias in classical generative vision models [15]. In both configurations, a two-hidden-layer MLP (128 and 256 units) maps the observable and positional vectors to one grayscale patch. Figure 1 shows the exact HVK1D and HVK2D ansatz circuits (rebuilt in Qiskit and numerically verified against the training-time PennyLane circuit as part of the hardware pilot of Section IV-D); complete Hamiltonian, loss, positional-encoding, and measurement definitions are retained in the companion supplement.

Table I summarizes the tested variants. HVK1D uses a linear 6-qubit chain (5 nearest-neighbor bonds, 27-D observable vector); HVK2D arranges the same 6 qubits on a 2×3 lattice matching patch-boundary adjacency (19-D latent signature); a symmetric HVK1D variant adds a $U(1)$ -preserving coupling constraint; and a D_4 -pooled HVK2D variant (Section IV-G) group-averages the observable grid over the eight symmetries of the square, $\Phi_{D_4}(x) = \frac{1}{8} \sum_{g \in D_4} P_g^{-1} \Phi(gx)$, which is equivariant by construction. The IBM Cloud circuit probe reported in the supplement confirms both HVK1D and HVK2D compile to depth-18, 6-qubit circuits on a Heron-class basis.

III. EXPERIMENTAL SETUP

All image-reconstruction results in Sections IV-A–IV-D use the single-image *Monalisa* protocol developed for HVK: a fresh model is optimized directly on its target image, with no

TABLE I
HVK VARIANTS TESTED IN THIS PAPER.

Variant	Topology	Latent dim.	Symmetry
HVK1D	6-qubit chain	27-D	None (control)
HVK2D	2×3 grid	19-D	D_4 -motivated only
Symmetric HVK1D	6-qubit chain	27-D	$U(1)$ coupling
D_4 -pooled HVK2D	2×3 grid, pooled	19-D	D_4 -equivariant (exact)

TABLE II
HVK2D SAME-SET RECONSTRUCTION ACROSS SIX REAL IMAGE DATASETS (PER-IMAGE TRAINED, 80–100 STEPS, MEAN $\pm$ STD OVER SAMPLED IMAGES).

Dataset	n images	PSNR (dB)	SSIM
CIFAR-10	3	37.75 ± 2.38	0.9986 ± 0.0003
MNIST	3	25.77 ± 2.23	0.9783 ± 0.0116
Fashion-MNIST	3	34.55 ± 3.05	0.9980 ± 0.0006
PathMNIST	3	36.86 ± 2.11	0.9118 ± 0.0627
BloodMNIST	2	36.57 ± 2.36	0.9922 ± 0.0063
PneumoniaMNIST	2	41.60 ± 0.98	0.9985 ± 0.0006

held-out split. They therefore measure per-image fitting and reconstruction capability, including hardware replay of those fitted checkpoints. The symmetry and restricted-entanglement diagnostics use the separate protocols stated in their respective sections; the exploratory training-dynamics material is addressed only as a provenance audit and is not retained as evidence. Held-out generalization and component necessity are evaluated under the resource-matched statistical protocol of the companion supplementary study (Section IV-C). All settings report MSE and $\text{PSNR} = 20 \log_{10}(1/\sqrt{\text{MSE}})$; image settings additionally report SSIM. Full implementation settings and environment versions are given in the companion supplementary document.

IV. RESULTS

A. Reconstruction Across Six Real Image Datasets

For each image, we trained a fresh HVK2D model using nonoverlapping 8×8 patches (16 patches/image, 80–100 optimization steps). The sample covers CIFAR-10 [32], MNIST [33], Fashion-MNIST [34], and the PathMNIST, BloodMNIST, and PneumoniaMNIST subsets of MedMNIST v2 [35]. Table II reports mean PSNR between 25.8 and 41.6 dB (SSIM 0.91–1.00). This is a same-set fitting result across varied image structures, not a held-out generalization claim; the companion study provides the latter evaluation (Section IV-C).

B. Scope and Adaptation Across Images

A model optimized on one image alone does not transfer zero-shot to a structurally different image: on *Monalisa*, zero-shot transfer to a second image reaches only 8.31 dB PSNR (SSIM = 0.1044), while extending training to include the second image recovers 28.63 dB (SSIM = 0.9858). This confirms that the per-image results above measure per-image optimization and reconstruction capability, not held-out generalization — consistent with how we scope every claim in this paper, and precisely the distinction the held-out protocol of Section IV-C is designed to test.




Fig. 1. The HVK1D (top) and HVK2D (bottom) variational ansätze, exactly as executed (encoding R_X rotations, positional R_Y rotations, then two entangling layers of R_Z - R_Y - R_Z single-qubit rotations with a CNOT ring for HVK1D or fixed grid-edge CNOTs for HVK2D). Both circuits were rebuilt gate-for-gate in Qiskit and validated against the original PennyLane training circuit before use in the hardware pilot of Section IV-D.

C. Competitive with Resource-Matched Classical Baselines

The results above measure per-image fitting; a separate, equally important question is whether HVK’s feature map improves reconstruction on real, *held-out* images under a resource-matched multi-seed protocol. That question is answered in the companion document, *Supplementary Material for the Hamiltonian Vision Kernel: Resource-Matched Ablations, Validation, and Reproducibility*. On held-out CIFAR-10, local-observable and raw-linear controls reach 18.80 ± 1.42 dB against HVK2D’s 18.12 ± 1.54 dB; after aggregating held-out images within each split, the HVK2D-minus-control difference is -0.68 dB (bootstrap 95% CI $[-0.91, -0.46]$ dB; Wilcoxon $p = 0.0625$, $n = 5$ seeds), with the same numerical direction on five further real datasets. Thus the tested HVK map is informative but is not shown to outperform the simpler maps on ordinary held-out reconstruction. A TOST equivalence test at a declared ± 1 dB margin (companion supplement) confirms this gap is small enough, in a pre-declared and data-independent sense, to call the two *competitive* rather than merely “not shown to differ”: six of eight tested controls are statistically equivalent to HVK2D within that margin, while the two weakest controls (strict classical random features, random-VQC) are correctly excluded because HVK2D beats them by a wide, practically meaningful margin. We report this as a rigorous complement: it delineates the tested boundary of the positive entanglement and symmetry results reported below and records the provenance checks that prevented exploratory artifacts from becoming claims. A matched-budget, multi-seed rerun of the component-ablation table (240 steps, 3 seeds/variant, replacing an earlier table whose per-seed provenance could not be recovered) reaches the same conclusion on ordinary per-image reconstruction: a plain classical-replacement control statistically ties or slightly exceeds the quantum baseline, while only a circuit-free random-VQC con-

trol performs clearly worse; full results are in the companion supplement. Section V-B discusses why this ties-with-classical result and the entanglement-necessity result of Section IV-H are not in tension.

D. A Real Hardware Reconstruction Pilot

Every result above is a noiseless-simulator result, a limitation we flagged rather than left implicit. We report a first step past it: a small but complete image-reconstruction pilot executed on real IBM Quantum hardware (`ibm_fez`, a 156-qubit Heron-class processor), beyond the compiled-circuit feasibility probe documented in the supplement. For five images — the full 16-patch *Monalisa* benchmark and four native-resolution CIFAR-10 images (*cat*, *ship* $\times 2$, *airplane*), each optimized directly on its own target image following the same per-image protocol as every result in this paper — we took an already-trained HVK1D/HVK2D checkpoint’s exact learned parameters, rebuilt its measurement circuits gate-for-gate in Qiskit, verified the translation against the original PennyLane simulator to within numerical noise ($\leq 7.9 \times 10^{-4}$ maximum absolute error per observable across every patch; full validation numbers in the companion supplementary document), and then executed the same circuits on real hardware with 256 shots per measurement basis, a fixed per-basis shot budget rather than an adaptive shadow-tomography-style estimation scheme [25]. The hardware-measured observables were decoded by the same trained classical decoder, with no retraining and no simulator involved anywhere in the reconstruction path.

Reconstructions from real, noisy hardware measurements retain recognizable image structure. *Monalisa* reaches 25.90 dB PSNR from hardware measurements, against a 37.99 dB noiseless-simulator replay of the identical circuit (Figure 2); the four CIFAR-10 images average 29.56 ± 2.03 dB on hardware against 42.69 ± 3.84 dB in simulation (Figure 3,

TABLE III
REAL IBM HARDWARE RECONSTRUCTION PILOT (BACKEND `ibm_fez`,
256 SHOTS/BASIS, NO RETRAINING). PSNR COMPARES
HARDWARE-DECODED AND SIMULATOR-DECODED RECONSTRUCTIONS
AGAINST THE SAME GROUND-TRUTH IMAGE.

Image	Patches	Sim. PSNR (dB)	HW PSNR (dB)
Monalisa (HVK1D)	16	37.99	25.90
CIFAR cat (HVK2D)	16	44.85	31.52
CIFAR ship, hydrofoil (HVK2D)	16	36.56	26.44
CIFAR ship, sea boat (HVK2D)	16	42.57	31.20
CIFAR airplane (HVK2D)	16	46.79	29.10
CIFAR mean $\pm$ std (4 images)	–	42.69 $\pm$ 3.84	29.56 $\pm$ 2.03

Table III). The observed gap is consistent with noise in these unmitigated device executions, although the five-image pilot does not isolate individual noise sources. Every hardware reconstruction remains above the random-latent control range (11–15 dB in the companion supplementary study). The pilot used 176 circuits; the contemporaneous account-meter reading increased by 16 quantum-processor seconds, with the provenance limitation documented in the supplement.

Five images constitute a pilot, and we make no hardware-advantage claim. The narrower result is that the trained HVK observable channel remains visually and quantitatively informative in these device executions under the reported circuit and shot budget. Full methodology — exact gate sequences, per-basis measurement construction, job IDs, and pre-submission local-validation results — is documented in the companion supplementary document.

The companion Supplementary Material reports a further hardware-validation exercise built on this same checkpoint-replay methodology: fresh HVK1D checkpoints trained at $N \in \{4, 6, 8\}$, replayed epoch-by-epoch on real IBM Quantum hardware and on the IonQ cloud simulator [26], in support of the training-dynamics diagnostic discussed in Section V-A and the Supplementary Material. We keep it out of the main text because it is a single-patch, single-seed probe in service of a diagnostic whose own negative control (below) does not establish sensitivity to the Hamiltonian term it is meant to track.

E. Noise/Shot Trade-off: A Free Simulator Complement

The five-image pilot above used a single shot budget (256/basis) and one execution per image, which bounds how much it can say about noise sensitivity on its own. Real IBM Quantum time is a scarce free-tier resource, so we separate what a *calibrated device-noise simulation* can establish at zero hardware cost from what genuinely requires hardware. Using `qiskit-aer` with the live calibration snapshot of `ibm_fez` (FakeFcz, no quantum-processor time consumed), we replay the same already-trained checkpoints’ circuits at four shot budgets (256/512/1024/4096 shots per basis, three repeated executions each) and compare against both the noiseless-simulator ceiling and the real-hardware point already reported above.

Two findings follow, averaged over the five checkpoints. First, shot noise accounts for only part of the noiseless-to-real gap: moving from the noiseless ceiling to a calibrated-

TABLE IV
REAL-HARDWARE ANCHOR POINTS: SAME CHECKPOINTS, NO
RETRAINING, ON TWO FURTHER BACKENDS. THE ORIGINAL PILOT’S
`ibm_fez` POINT IS SHOWN FOR REFERENCE.

Image	Backend	Shots	PSNR (dB)
Monalisa	<code>ibm_fez</code> (original pilot)	256	25.90
Monalisa	<code>ibm_marrakesh</code>	256	25.94
Monalisa	<code>ibm_kingston</code>	1024	26.10
CIFAR cat	<code>ibm_fez</code> (original pilot)	256	31.52
CIFAR cat	<code>ibm_kingston</code>	256	28.93
CIFAR cat	<code>ibm_kingston</code>	1024	31.24

noise simulation at the same 256-shot budget already costs 8.68 dB on average, before any real device is involved. Second, more shots does not reliably recover it: increasing to 4096 shots per basis changes PSNR by only -0.05 dB on average across the five checkpoints, with individual images ranging from a 1.33 dB gain to a 0.96 dB loss — at this circuit size and shot range, shot count is not a lever that reliably improves reconstruction quality. Comparing the 256-shot noisy simulation directly against the real-hardware point isolates what the calibration snapshot does *not* capture: a further 4.24 dB average gap (range 2.25–8.06 dB across images), attributable to device effects a single calibration snapshot does not fully capture — drift between calibration and execution, crosstalk, and correlated errors chief among them. The airplane image is the clear outlier (8.06 dB residual gap, roughly double any other image), which we report rather than exclude; a single-execution, five-image pilot cannot distinguish an image-specific effect from ordinary job-to-job variation. This calibrated-noise simulation is a free, repeatable complement to the pilot, not a substitute for it: it bounds how much of the observed hardware gap is attributable to modeled device noise versus unmodeled effects, using zero additional quantum-processor time.

F. Real-Hardware Anchor Points: Repeated Execution and a Second Backend

The simulated trade-off above is bounded, in turn, by whether it actually reflects real devices. We spent a small, explicitly quota-tracked slice of the free-tier IBM Quantum allowance (four jobs, well under 1% of the monthly limit) confirming it directly: the same two checkpoints (*Monalisa*, *CIFAR cat*) replayed with no retraining on `ibm_marrakesh` and `ibm_kingston` — two backends distinct from the original pilot’s `ibm_fez` — at the original 256-shot budget and, separately, at 1024 shots.

Monalisa reproduces almost exactly across backends at matched shots (25.94 vs. 25.90 dB, `ibm_marrakesh` vs. `ibm_fez`), and gains only 0.16 dB from quadrupling shots to 1024 — consistent with the simulated finding above that shot count is not a reliable lever. The *CIFAR cat* point is the more informative one: on `ibm_kingston` at the original 256-shot budget it reaches 28.93 dB, 2.59 dB below the `ibm_fez` pilot point, while its own 1024-shot execution on the same backend recovers to 31.24 dB, within 0.28 dB of the original. This is real job-to-job and backend-to-backend variation of a magnitude comparable to the unmodeled residual



Fig. 2. *MonaLisa* reconstructed from real IBM Quantum hardware measurements (right) against the same trained checkpoint’s noiseless-simulator replay (center) and ground truth (left). The hardware image is visibly noisier and exhibits patch-boundary banding, while the subject remains recognizable.

TABLE V
 D_4 EQUIVARIANCE ERROR FOR UNPOOLED AND POOLED OBSERVABLE-GRID MAPS (LOWER IS BETTER; POOLED MAP IS GROUP-AVERAGED OVER ALL EIGHT SQUARE SYMMETRIES).

Model	Mean error	Std. error
D_4 -pooled HVK2D	9.57×10^{-17}	1.26×10^{-17}
No positional encoding	7.40×10^{-1}	1.52×10^{-1}
HVK2D positional	8.15×10^{-1}	1.29×10^{-1}
Local/raw	8.37×10^{-1}	1.31×10^{-1}

identified by the simulator comparison (Section IV-E) — direct evidence, not merely a simulated bound, that a single 256-shot execution on a single backend is not yet a stable estimate of this pipeline’s hardware performance, and that repeated execution matters at least as much as backend choice or shot count individually.

G. Exact D_4 -Equivariant Observable Pooling

A D_4 -pooled HVK2D observable-grid map is equivariant to the eight symmetries of the square *by construction*: $\Phi_{D_4}(x) = \frac{1}{8} \sum_{g \in D_4} P_g^{-1} \Phi(gx)$ averages the unpooled map over the group orbit, so $\Phi_{D_4}(hx) = P_h \Phi_{D_4}(x)$ for every $h \in D_4$ up to floating-point error. We verify this numerically over 7000 patch transforms of cached CIFAR images (Table V, Figure 5): the pooled map’s equivariance error is $9.57 \times 10^{-17} \pm 1.26 \times 10^{-17}$ (floating-point noise), while the original unpooled positional map is not measurably more symmetric than a local/raw control (0.815 ± 0.129 vs. 0.837 ± 0.131) — confirming the unpooled map is only D_4 -motivated by the patch grid’s geometry, not empirically equivariant, and that only the explicit pooling construction delivers the property. This is currently a feature-grid-level construction and numerical diagnostic rather than a component of the trainable reconstruction pipeline; integrating it end-to-end (Section V-F) is concrete future work.

H. Entanglement Necessity in the Nonlocal-Structure Protocol

This diagnostic is deliberately constructed to reward pair-observable structure: the target is a fixed synthetic function of *simulated* patch features not concatenated into any candidate’s own input block (the companion supplement gives the leakage-audit procedure that rules out circularity), and every control receives the same raw inputs as HVK2D. Table VI and Figure 6 report the result across five seeds: HVK2D’s entangling observable channel reaches mean $R^2 = 0.9735$ (32.77 dB PSNR), a 17.74 dB improvement over a no-entanglement control ($R^2 = 0.019$) and 18.52 dB over random-VQC latents ($R^2 = -0.173$); freezing the quantum circuit (features fixed, only readout trained) still reaches $R^2 = 0.853$, while a parameter-matched classical control ($R^2 = -0.030$) and a frozen-classical control ($R^2 = -0.903$) fail outright. This is evidence that an entangling pair-observable channel is *representationally necessary within the tested feature class and fixed linear-readout protocol* to fit this target.

I. A Scoped Training-Dynamics Diagnostic (Full Account in the Supplementary Material)

Beyond the results above, we also tracked two exploratory training-dynamics diagnostics: a magnetization-style change-point statistic $M_z(t)$ and a signed energy-to-entanglement ratio $R_{ES}(t)$, varied along qubit-count ($N = 4, 6, 8$) and MPS-bond-dimension ($\chi = 1, 2, 4, 8$) axes and partly replayed on real IBM Quantum and IonQ hardware. We summarize the outcome here and give the full protocol, tables, and figures in the companion Supplementary Material, because the diagnostic does not clear the bar for a main-text result: applied to a full six-dataset, two-seed sample, it detects a change in 16/24 runs, a mixed and non-universal rate; and its own negative control — the identical protocol applied to a classical-replacement variant with the Hamiltonian energy identically zero throughout training — fires in all 4/4 tested cases, exactly as often as the Hamiltonian-regularized model’s 4/4. The detector therefore does not distinguish a Hamiltonian-driven change from one this descriptive, per-run threshold would find




Fig. 3. Four CIFAR-10 images reconstructed from real IBM Quantum hardware measurements (bottom row) against noiseless-simulator replays of the same trained HVK2D checkpoints (middle row) and ground truth (top row).

TABLE VI
HVK2D RESTRICTED PAIR-CORRELATION DIAGNOSTIC ACROSS FIVE SEEDS. DELIBERATELY FAVORABLE TO PAIR-OBSERVABLE STRUCTURE; EVIDENCE FOR ENTANGLEMENT-NECESSITY UNDER A MATCHED LINEAR READOUT.

Model	Mean MSE	Mean PSNR (dB)	Mean R^2
HVK2D entangling observables	8.5969×10^{-4}	32.77	0.9735
Freeze quantum	4.7904×10^{-3}	27.34	0.8533
No entanglement	3.1461×10^{-2}	15.03	0.0191
Raw-linear classical	3.1654×10^{-2}	15.00	0.0133
Parameter-matched classical	3.3033×10^{-2}	14.82	-0.0297
Random VQC	3.7616×10^{-2}	14.25	-0.1732
Freeze classical	6.0973×10^{-2}	12.17	-0.9025

regardless of whether any physical energy term is present at all. We read this as a diagnostic whose own control undercuts it, not as evidence of a physical transition, and report it at that scope rather than as a headline result. The system-size and bond-dimension sensitivity checks in the supplement reinforce the same reading: the detected change epoch is non-monotonic in qubit count under one aggregation and monotonic under another, and is comparably variable across all tested MPS bond dimensions — instability of the heuristic, not a scaling law.

V. DISCUSSION

A. Why the Physics-Informed Regularizer Was Inert

Table VII shows the Hamiltonian energy term is, if anything, a mild drag on reconstruction quality in the present system. These are fresh, same-environment measurements (*Monalisa*, 200 steps, seed 42); an earlier draft reported different absolute PSNR values for these same three configurations, produced by a script that was never committed to the repository and cannot be recovered or audited, so those earlier numbers are superseded here rather than cited (full account of the discrep-




Fig. 4. Calibrated device-noise simulation (FakeFez, mean $\pm$ std over three executions). Dotted lines show noiseless ceilings and stars show real-hardware measurements. Increasing 256 to 4096 shots changes mean PSNR by only -0.05 dB.




Fig. 5. D_4 equivariance diagnostic. The unpooled positional map is not equivariant; the group-pooled map is equivariant to numerical precision.

ancy, plus additional observable-sector controls and follow-up experiments, in the Supplementary Material). The legacy objective adds a signed mean energy to the reconstruction loss; because learned energy can go negative, the optimizer can lower total loss by driving energy down without improving reconstruction, and removing the term improves PSNR from 38.63 to 44.56 dB. A contrastive variant — assigning lower energy to correct feature–position pairs than to shuffled pairs — improves over the signed-energy baseline (42.92 dB) but remains below its own no-energy-loss counterpart (44.56 dB). We read this as a case where the reconstruction loss already dominates the optimization signal at this model scale and dataset size: a physically motivated regularizer has room to help when it resolves an otherwise underdetermined objective, and here the pixel-fidelity loss alone is sufficiently informative that the extra term mostly adds noise to the gradient. This

does not mean Hamiltonian regularization is never useful — only that its value is conditional on the reconstruction loss being under-constrained, a condition this patch-level pixel task does not meet. The Supplementary Material additionally reports a root-cause fix (bounding the couplings and using a positive/squared energy loss, which stops the term from hurting but does not make it help) and an exploratory decoder-feature variant that improves 2 of 3 seeds but regresses sharply on the third, too seed-sensitive to report as a validated result here.

The order-parameter and energy-to-entanglement-ratio diagnostics of Section IV-I and the companion Supplementary Material provide descriptive windows into training dynamics. Their detected epochs come from changes in measured model summaries, but the data-dependent threshold is not a calibrated statistical test and the results should not be interpreted as thermodynamic transitions.

B. Task-Dependent Representational Scope

Whether a quantum feature map or kernel offers any advantage over a classical model is a property of the specific data and target, not a fixed property of the model class, and dequantization results have repeatedly shown that apparent quantum advantages on structured, low-entanglement classical data can be matched by classical algorithms once the right classical feature is identified [27]. The restricted diagnostic of Section IV-H is constructed so that the relevant structure is a genuine distant-element product raw or local features cannot represent under a fixed linear readout, and that is precisely the regime in which HVK’s entangling channel separates. Natural image patches, at the resolution and patch size used in ordinary reconstruction, are a different regime — low-entanglement, locally structured data where a linear or shallow classical readout is already close to sufficient. The companion supplementary study’s held-out comparisons (Section IV-C) characterize exactly this second regime; together, the results delineate where the tested entangling representation does and does not yield a measurable benefit.

C. HVK as a Representative Instance of a Broader Design Pattern

The boundary reported in Section IV-C raises an immediate question: is it a property of HVK’s particular implementation choices, or of the broader class of architectures HVK instantiates? HVK combines three ingredients that recur, individually, throughout the hybrid quantum-vision literature rather than being idiosyncratic to this codebase: matrix-product-state compression of structured input into a classical or hybrid feature vector [7], [9], [11], [22]; a Pauli-correlator (measured-observable) latent space handed to a classical decoder rather than a raw quantum state, in the spirit of quantum-enhanced feature maps [21]; and a physically motivated auxiliary regularizer alongside a separate task loss [19]. HVK is not a minimal or contrived combination of these three choices assembled to produce a particular result; each is independently




Fig. 6. Restricted pair-correlation diagnostic summary (mean $\pm$ std over five seeds). The entangling observable feature map dominates every control on this deliberately favorable task.

TABLE VII

HAMILTONIAN-SECTOR CONTROLS (SINGLE-SEED, FRESH SAME-ENVIRONMENT REPRODUCTION: *Monalisa*, 200 STEPS, SEED 42). THE ENERGY OBJECTIVE IS DIAGNOSTICALLY USEFUL BUT DOES NOT IMPROVE RECONSTRUCTION ACCURACY. SEE THE SUPPLEMENTARY MATERIAL FOR THE FULL REPRODUCIBILITY INVESTIGATION AND ADDITIONAL OBSERVABLE-SECTOR CONTROLS.

Experiment	PSNR (dB)	Interpretation
Legacy baseline, signed energy	38.63	Standard energy-regularized run
No energy loss	44.56	Removing energy improves PSNR
Contrastive Hamiltonian core	42.92	Stronger objective helps modestly vs. baseline, still below no-energy-loss

well-motivated and widely used, and HVK’s specific instantiation (six-qubit VQC, 19–27-D Pauli-correlator vector, graph-Hamiltonian energy term) sits well within the parameter ranges typical of this literature rather than at an unusual extreme.

This representativeness matters for ruling out an uncharitable alternative reading of the held-out result (Section IV-C) and the inert-regularizer result (Section V-A): that HVK’s specific implementation choices — an unmotivated qubit count, an idiosyncratic latent width, or an unusual regularizer formulation — rather than the design pattern itself, are responsible for tying classical baselines, and that a different, non-representative choice within the same pattern would have won. Because HVK’s instantiation sits well within the ordinary parameter ranges of this literature rather than at an unusual extreme, that alternative reading is not supported: nothing about HVK’s specific engineering choices explains away the result. Per the data-dependent framing of Section V-B [27], whether any instance of this design pattern outperforms a resource-matched classical alternative is a property of the specific data and target, not of the pattern’s architectural motivation, so we do not claim the measured boundary transfers unchanged to every dataset, patch size, or qubit count. What this section establishes is narrower and, we think, more useful: the architecture class HVK instantiates — MPS feature compression, a Pauli-correlator latent space, and a physics-informed regularizer — is sound and general, a fair and properly-scaled test of the pattern rather than a strawman version of it, so its measured behavior here

is evidence worth taking at face value rather than dismissing as an idiosyncrasy of one implementation.

D. Topology Alignment as an Inductive Bias

Aligning the latent interaction graph with image-patch adjacency (HVK2D vs. HVK1D) is a plausible architectural inductive bias. The companion supplementary study tests this hypothesis directly under matched held-out conditions. Its real-circuit subset (3 seeds, 3 held-out CIFAR-10 images) gives an HVK1D-minus-HVK2D mean PSNR difference of 0.16 dB; after averaging images within seed, the bootstrap 95% CI is $[-0.38, 0.46]$ dB and the seed-level Wilcoxon test gives $p = 0.50$. HVK2D completes the matched 90-step runs in roughly half the wall-clock time. A broader five-seed surrogate comparison across CIFAR-10, Fashion-MNIST, and PathMNIST likewise finds numerically small topology differences (at most 0.07 dB in mean PSNR). At the tested scale, topology therefore appears to be primarily a resource-cost choice rather than a reliable accuracy advantage; these tests do not establish formal equivalence.

E. Architectural Differentiation Matrix

Table VIII contrasts HVK’s structural properties against reference frameworks, including the three nearest hybrid-vision architecture families discussed above (register-based QAE, quantum-kernel vision, TN-circuit classifier). Its distinguishing features — a Pauli-correlator latent space, 2D

Fourier positional bias with enforced D_4 pooling, and a Hamiltonian energy diagnostic — are documented and scoped in the preceding sections.

F. Validated Scope and Next Steps

- Every reconstruction result in this paper (Sections IV-A–IV-B) is per-image trained, not held-out generalization; the companion supplementary study’s held-out protocol is the appropriate evidence for a generalization claim, and its finding that local/raw maps match or exceed the tested HVK map under the shared readout should be read alongside the architecture-level results here.
- The D_4 -pooled map (Section IV-G) is currently a feature-grid-level construction and numerical diagnostic, not yet integrated into the trainable end-to-end reconstruction pipeline; doing so is concrete future work (Section VI).
- The entanglement-necessity result (Section IV-H) is scoped to a task deliberately constructed to require nonlocal structure; it does not by itself establish an advantage on ordinary natural-image reconstruction (see Section V-B).
- The hardware pilot (Section IV-D) is a five-image, single-trial, unmitigated pilot, not a statistically resolved hardware benchmark.
- The original exploratory training-dynamics traces were withdrawn after a provenance audit (Section IV-I); the corrected, noise-free multi-seed replacement is a small-sample diagnostic (2 images $\times$ 2 seeds/dataset, 24 runs across six datasets) with mixed, not universal, detection (16/24), and should be read at that scope.
- The change-point detector’s negative control (Section IV-I, full table in the Supplementary Material) fires on a classical-replacement variant with the Hamiltonian energy identically zero exactly as often as on the Hamiltonian-regularized model (4/4 vs. 4/4). The diagnostic should therefore be read as a property of the descriptive threshold applied to the change-magnitude trace, with no demonstrated sensitivity to the Hamiltonian term itself.
- The matched real-circuit topology comparison uses only 3 seeds and 3 held-out images per topology; the broader five-seed comparison is surrogate-based. These results support a resource-cost interpretation but do not establish formal equivalence of the topologies.

VI. CONCLUSION

We presented the Hamiltonian Vision Kernel, a new hybrid quantum–classical architecture for image reconstruction that assembles matrix-product-state patch features, a measured Pauli-correlator latent space, enforced D_4 -equivariant observable pooling, and a learnable Hamiltonian energy diagnostic into a single pipeline — a combination not previously assembled in this form (Section V-C). Across six real image datasets, per-image-trained HVK2D models reconstruct with high fidelity, and on a resource-matched, multi-seed held-out comparison a pre-declared TOST equivalence test shows HVK2D is *competitive with*, rather than inferior or superior to,

simple classical controls (Section IV-C): local/raw maps match or exceed it under the same fitted readout, and TOST confirms the gap is small enough to call the two practically tied. We further reported a real IBM Quantum hardware reconstruction pilot: replaying trained checkpoints without retraining yielded 25.90–31.52 dB PSNR from hardware measurements, establishing feasibility at a five-image, single-trial scope rather than a general hardware-performance claim.

If HVK ties classical baselines on ordinary reconstruction, its case rests on the capabilities a purely classical reconstruction pipeline does not provide: a feature-level, numerically verified D_4 -equivariance construction (equivariance error 9.57×10^{-17}), and a restricted entanglement-sensitive result on a task requiring nonlocal structure, where pair observables are necessary within the tested fixed-readout feature class ($R^2 = 0.9735$ vs. $R^2 \leq 0.02$ for non-entangling controls). Neither capability is in tension with the tie-with-classical result: entangling observables separate only on the tested nonlocal target, and group pooling guarantees equivariance without by itself establishing an accuracy gain. We make no claim of quantum advantage. HVK is, at this scope, a new architecture that works, is competitive with resource-matched classical baselines, runs on real quantum hardware, and is equipped with an entanglement-sensitive channel, an interpretable Hamiltonian diagnostic, and native hardware execution — capabilities a classical reconstruction pipeline does not natively expose. We consider this carefully measured boundary a constructive contribution alongside the architecture itself.

Future work should integrate the D_4 -pooled map into an end-to-end trainable pipeline (currently a feature-grid-level construction, Section V-F); scale the hardware pilot toward a statistically resolved benchmark with error mitigation and repeated trials; search for further real-world tasks with genuine nonlocal structure where HVK’s restricted entanglement-sensitive separation translates into a reconstruction advantage; and extend the resource-matched, leakage-audited evaluation protocol used in the companion supplementary study to other hybrid quantum vision architectures.

DATA AND CODE AVAILABILITY

The source code, experiment drivers, machine-readable result summaries, validation arrays, and figure inputs supporting this study are included in the accompanying repository. The file `project_artifacts/submission_claim_audit.json` maps the three numerical headline claims to their retained source artifacts, and automated tests verify those mappings. Dataset provenance and access routes are cited in the manuscript and detailed in the supplementary artifact map. IBM Quantum job identifiers and reconstructed outputs are retained; as disclosed in the supplement, the raw account-usage service response was not serialized, so the reported quota change is an execution-log record rather than a repository-recomputable measurement.

ACKNOWLEDGEMENTS

We thank the maintainers of the open-source scientific computing libraries used in this work.

TABLE VIII

ARCHITECTURAL DIFFERENTIATION FROM REPRESENTATIVE BASELINE DESIGNS, INCLUDING THE THREE NEAREST HYBRID-VISION ARCHITECTURE FAMILIES (REGISTER-BASED QAE, QUANTUM-KERNEL VISION, TN-CIRCUIT CLASSIFIER).

Axis	Conv. AE baseline	GAN baseline	Register-based QAE	Quantum-kernel vision	TN-circuit classifier	HVK (ours)
Latent representation	Continuous $\mathbb{R}^d$	Continuous $\mathbb{R}^d$	Quantum register	Scalar kernel value	MPS/circuit state	Pauli-correlator vector
Spatial bias	Convolutional	Convolutional	Encoding-dependent	Encoding-dependent	MPS chain ordering	2D Fourier encoding
Auxiliary objective	Bottleneck/reconstruction	Adversarial loss	Fidelity/reconstruction	None (kernel value only)	Classification loss	Hamiltonian diagnostic
Enforced group symmetry	Not in baseline	Not in baseline	Not in reference design	Not in reference design	Not in reference design	D_4 observable pooling
Evaluated role	Reconstruction	Reconstruction	Compression	Classification	Classification	Reconstruction + diagnostics

REFERENCES

- [1] M. Cerezo et al., “Variational quantum algorithms,” *Nature Reviews Physics*, vol. 3, no. 9, pp. 625–644, 2021.
- [2] J. Tilly et al., “The Variational Quantum Eigensolver: A review of methods and best practices,” *Physics Reports*, vol. 986, pp. 1–128, 2022.
- [3] B. Bauer, S. Bravyi, M. Motta, and G. K.-L. Chan, “Quantum Algorithms for Quantum Chemistry and Quantum Materials Science,” *Chemical Reviews*, vol. 120, no. 22, pp. 12685–12717, 2020.
- [4] S. Chakraborty, S. B. Mandal, and S. H. Shaikh, “Quantum image processing: challenges and future research issues,” *International Journal of Information Technology*, vol. 14, no. 1, pp. 475–489, 2022, doi: 10.1007/s41870-018-0227-8.
- [5] R. Zhou, W. Hu, P. Fan, and H. Ian, “Quantum realization of the bilinear interpolation method for NEQR,” *Scientific Reports*, vol. 7, no. 1, p. 2511, 2017.
- [6] X. Yao et al., “Quantum Image Processing and Its Application to Edge Detection: Theory and Experiment,” *Physical Review X*, vol. 7, no. 3, p. 031041, 2017.
- [7] J. I. Cirac, D. Pérez-García, N. Schuch, and F. Verstraete, “Matrix product states and projected entangled pair states: Concepts, symmetries, theorems,” *Reviews of Modern Physics*, vol. 93, no. 4, p. 045003, 2021.
- [8] J. I. Latorre, “Image compression and entanglement,” arXiv:quant-ph/0510031, 2005.
- [9] Z. Han, J. Wang, H. Fan, L. Wang, and P. Zhang, “Unsupervised Generative Modeling Using Matrix Product States,” *Physical Review X*, vol. 8, no. 3, p. 031012, 2018.
- [10] S. M. Fei, “Compressed Sensing Based on Tensor Network Machine Learning,” *Physical Science & Biophysics Journal*, vol. 5, no. 1, Art. no. 000174, 2021, doi: 10.23880/psbj-16000174.
- [11] D. Guala et al., “Practical overview of image classification with tensor-network quantum circuits,” *Scientific Reports*, vol. 13, no. 1, p. 4427, 2023.
- [12] A. Haga, “Quantum annealing-based computed tomography using variational approach for a real-number image reconstruction,” *Physics in Medicine and Biology*, vol. 69, no. 4, Art. no. 04NT02, 2024, doi: 10.1088/1361-6560/ad2155.
- [13] X. Zhai, T. Xiao, J. Huang, J. Fan, and G. Zeng, “Quantum neural compressive sensing for ghost imaging,” *Physical Review Applied*, vol. 23, no. 1, p. 014018, 2025.
- [14] M. A. Nau, A. H. Vija, W. Gohn, M. P. Reymann, and A. Maier, “Exploring the Limitations of Hybrid Adiabatic Quantum Computing for Emission Tomography Reconstruction,” *Journal of Imaging*, vol. 9, no. 10, p. 221, 2023.
- [15] R. Xu, X. Wang, K. Chen, B. Zhou, and C. C. Loy, “Positional Encoding as Spatial Inductive Bias in GANs,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021, pp. 13564–13573.
- [16] J. Y. Araz and M. Spannowsky, “Classical versus quantum: Comparing tensor-network-based quantum circuits on Large Hadron Collider data,” *Physical Review A*, vol. 106, no. 6, p. 062423, 2022.
- [17] D. Slabbert and F. Petruccione, “Hybrid Quantum-Classical Feature Extraction approach for Image Classification using Autoencoders and Quantum SVMs,” arXiv:2410.18814, 2024.
- [18] Z. Liu, P.-X. Shen, W. Li, L.-M. Duan, and D.-L. Deng, “Quantum capsule networks,” *Quantum Science and Technology*, vol. 8, no. 1, p. 015016, 2022.
- [19] J. Shi, W. Wang, X. Lou, S. Zhang, and X. Li, “Parameterized Hamiltonian Learning With Quantum Circuit,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 45, no. 5, pp. 6086–6095, 2022.
- [20] J. Preskill, “Quantum Computing in the NISQ era and beyond,” *Quantum*, vol. 2, p. 79, 2018.
- [21] V. Havlíček, A. D. Córcoles, K. Temme, A. W. Harrow, A. Kandala, J. M. Chow, and J. M. Gambetta, “Supervised learning with quantum-enhanced feature spaces,” *Nature*, vol. 567, no. 7747, pp. 209–212, 2019.
- [22] E. M. Stoudenmire and D. J. Schwab, “Supervised Learning with Tensor Networks,” in *Advances in Neural Information Processing Systems*, vol. 29, 2016.
- [23] T. Cohen and M. Welling, “Group Equivariant Convolutional Networks,” in *Proceedings of the 33rd International Conference on Machine Learning*, PMLR, vol. 48, 2016, pp. 2990–2999.
- [24] J. J. Meyer, M. Mularski, E. Gil-Fuster, A. A. Mele, F. Arzani, A. Wilms, and J. Eisert, “Exploiting Symmetry in Variational Quantum Machine Learning,” *PRX Quantum*, vol. 4, p. 010328, 2023.
- [25] H.-Y. Huang, R. Kueng, and J. Preskill, “Predicting many properties of a quantum system from very few measurements,” *Nature Physics*, vol. 16, pp. 1050–1057, 2020.
- [26] IonQ, “Backends,” *IonQ Quantum Cloud Documentation*. [Online]. Available: <https://docs.ionq.com/user-manual/backends>. Accessed: Jul. 24, 2026.
- [27] H.-Y. Huang, M. Broughton, M. Mohseni, R. Babbush, S. Boixo, H. Neven, and J. R. McClean, “Power of data in quantum machine learning,” *Nature Communications*, vol. 12, no. 1, p. 2631, 2021.
- [28] J. Romero, J. P. Olson, and A. Aspuru-Guzik, “Quantum autoencoders for efficient compression of quantum data,” *Quantum Science and Technology*, vol. 2, p. 045001, 2017. arXiv:1612.02806.
- [29] H. Wang, J. Tan, Y. Huang, and W. Zheng, “Quantum image compression with autoencoders based on parameterized quantum circuits,” *Quantum Information Processing*, vol. 23, p. 41, 2024.
- [30] H. Asaoka and K. Kudo, “Quantum autoencoders for image classification,” *Quantum Machine Intelligence*, vol. 7, p. 71, 2025. arXiv:2502.15254.
- [31] D. Beaulieu, D. Miracle, A. Pham, and W. Scherr, “Quantum Kernel for Image Classification of Real World Manufacturing Defects,” 2022. arXiv:2212.08693.
- [32] A. Krizhevsky, “Learning multiple layers of features from tiny images,” University of Toronto, Technical Report, 2009.
- [33] Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner, “Gradient-based learning applied to document recognition,” *Proceedings of the IEEE*, vol. 86, no. 11, pp. 2278–2324, 1998, doi: 10.1109/5.726791.
- [34] H. Xiao, K. Rasul, and R. Vollgraf, “Fashion-MNIST: A novel image dataset for benchmarking machine learning algorithms,” arXiv:1708.07747, 2017.
- [35] J. Yang et al., “MedMNIST v2—A large-scale lightweight benchmark for 2D and 3D biomedical image classification,” *Scientific Data*, vol. 10, Art. no. 41, 2023, doi: 10.1038/s41597-022-01721-8.