

HVK: A Hamiltonian-Guided Tensor Network Architecture for Spatially Aware Quantum Image Reconstruction

Sparsho Chakraborty* and Dr. Siddhartha Patra†

Abstract—We present the HVK architecture, an end-to-end quantum image reconstruction framework that integrates Matrix Product States (MPS), parameterized variational quantum circuits, and data-adaptive Heisenberg Hamiltonian operators. Classical approaches scale poorly with resolution, while typical quantum methods treat variational states opaquely without exploiting spatial connections. HVK addresses this by leveraging a local spatial smoothness regularizer explicitly derived from the expectation values of a low-energy spin-1/2 chain/grid configuration. We evaluate HVK on a Mona Lisa reconstruction task, a five-image CIFAR-10 subset, a reflection-constrained Symmetric HVK1D variant, and a small IBM Heron hardware probe. The results show that the 2D HVK topology improves over the 1D chain on reconstruction metrics and strongly outperforms GAN, autoencoder, and segmentation-oriented Hamiltonian baselines, while CNN and direct MLP baselines remain strong classical references. On IBM Heron, a 64-patch HVK2D probe produces stable positive order and nearest-neighbor ZZ correlations, validating the hardware measurability of the Hamiltonian latent signal.

Index Terms—Quantum Image Processing, Tensor Networks, Variational Quantum Algorithms, Heisenberg Hamiltonian, Image Reconstruction.

I. INTRODUCTION

Classical approaches—such as regularized inversion and, more recently, deep convolutional priors—have driven decades of progress, but their computational and memory costs scale poorly with image resolution and dimensionality. This has motivated a continuing search for representations and algorithms with more favorable scaling.

Quantum computing has emerged as one candidate paradigm. Variational quantum algorithms have demonstrated that parameterized quantum circuits, optimized in a hybrid quantum–classical loop, can prepare physically meaningful states on near-term hardware [1], [2]. The variational quantum eigensolver and its descendants—including the variational Hamiltonian ansatz and ADAPT-VQE—exploit the local decomposability of physically relevant Hamiltonians to construct shallow, problem-aware circuits whose parameters have transparent interpretations as rotation angles in a structured ansatz [1]–[3]. These developments raise a natural question: can Hamiltonian-guided variational principles, originally designed for electronic-structure problems, be repurposed for image

reconstruction, where the “Hamiltonian” encodes imaging physics rather than electronic interactions?

Quantum image processing (QIMP) provides the encoder side of such a pipeline. The flexible representation of quantum images encodes pixel colors in qubit rotation angles and pixel positions in computational-basis states, while the novel enhanced quantum representation (NEQR) stores grayscale values deterministically in basis states using q -qubit registers [4]. Successive works have improved and extended these encodings: quantum realizations of classical interpolation have been demonstrated for NEQR [5], and full quantum edge-detection pipelines for NEQR-encoded images have been validated theoretically and experimentally [6]. Quantum probability image encodings further reduce qubit and gate costs while preserving the structure needed for downstream processing [6]. However, most QIMP pipelines treat the encoded image as an opaque state to be acted upon; they do not expose the spatial correlations among neighboring pixels in a form that admits efficient variational optimization.

Tensor networks provide exactly this missing structure. Originally developed as compact representations of many-body quantum states, matrix product states (MPS), projected entangled pair states (PEPS), and the multi-scale entanglement renormalization ansatz (MERA) have proven remarkably effective as compressed representations of structured classical data, including images [7]. Classical tensor-network methods have been applied to image compression [8], to probabilistic generative modeling under partial observation [9], and—relevantly for inverse problems—to compressed sensing reconstructions whose latent state is itself a tensor network [10]. MPS/PEPS-structured quantum circuits have also been proposed as classical–quantum bridges that are qubit-efficient and well-suited to circuit-cutting techniques on near-term devices [11].

Despite this rich prior work, two streams remain largely disconnected. On one side, classical tensor-network formulations for imaging encode the image as an MPS and perform contraction and optimization without invoking quantum Hamiltonian dynamics; on the other, variational quantum algorithms for image reconstruction encode the image as a generic parameterized state without exposing the spatial bond structure of a tensor network [12]–[14]. The intersection—a Hamiltonian whose evolution respects and is explicitly coupled to a tensor-network topology that mirrors pixel adjacency—has not been systematically developed.

In this paper we introduce the HVK architecture, a

*School of Basic Sciences, Department of Physics, Indian Institute of Technology Bhubaneswar, Odisha, India.

† Assistant Professor, Centre for Quantum Engineering, Research and Education (CQUERE), TCG CREST, Kolkata, India.

Hamiltonian-guided tensor-network framework for spatially aware quantum image reconstruction. Rather than treating the image-encoding circuit and the reconstruction prior as separate modules, HVK fuses them: a problem-specific Hamiltonian is constructed such that its low-energy subspace, when probed via a variational ansatz organized as a tensor network, preferentially contains states consistent with both the measurement physics and the spatial structure of the scene.

II. THEORETICAL FRAMEWORK

A. Hamiltonian Operator Dynamics

HVK treats each spatial patch of an input image as a quantum many-body state on a finite spin-1/2 chain, encodes that state as a matrix product state, and exposes a Hamiltonian whose ground-state structure induces spatially aware correlations among neighboring pixels. This establishes the operator-theoretic foundations: the Heisenberg family of Hamiltonians, the corresponding energy objective, and the symmetry structure inherited by the resulting observable manifold.

1) *The Heisenberg Family of Hamiltonians:* For a one-dimensional chain of n_{bonds} bonds indexed by i and sites $(i, i+1)$, the anisotropic Heisenberg Hamiltonian is defined as:

$$H_{\text{Heis}} = \sum_{i=0}^{n_{\text{bonds}}-1} \left(J_x^{(i)} X_i X_{i+1} + J_y^{(i)} Y_i Y_{i+1} + J_z^{(i)} Z_i Z_{i+1} \right) \quad (1)$$

where X, Y, Z are the standard Pauli operators and $J_\alpha^{(i)}$ are bond-resolved coupling strengths. In the isotropic limit ($J_x = J_y = J_z$) the model possesses an exact $\text{SU}(2)$ symmetry; for $J_x = J_y \neq J_z$ (XXZ model) the $\text{U}(1)$ symmetry generated by $\sum Z_i$ is preserved, while for generic couplings (XYZ model) this symmetry is broken to $\mathbb{Z}_2 \times \mathbb{Z}_2$ [15]. In HVK, we fix the network context to six sites and treat every component $J_\alpha^{(i)}$ as a learnable real parameter, initialized from:

$$J_\alpha^{(i)} \sim \mathcal{N}(0, 0.01), \quad \alpha \in \{x, y, z\}. \quad (2)$$

Because the coupling strengths are optimized jointly with the circuit weights, the Hamiltonian itself adapts to image data, and the effective spin-spin interaction pattern is selected by the model dynamically.

2) *Energy Loss and Ground-State Objective:* Given a patch n and the Pauli observables measured from the variational circuit, the Heisenberg energy of the corresponding state is:

$$\mathcal{E}_n(\theta, J) = \sum_{i=0}^4 \sum_{\alpha \in \{x, y, z\}} J_\alpha^{(i)} \langle \sigma_\alpha^{(i)} \sigma_\alpha^{(i+1)} \rangle_n. \quad (3)$$

where α is the indices running from 1 to 3 denoting the Pauli matrices.

The training energy loss is evaluated as the patch-averaged energy:

$$\mathcal{L}_{\text{energy}}(\theta, J) = \langle E \rangle = \frac{1}{N} \sum_{n=1}^N \mathcal{E}_n(\theta, J). \quad (4)$$

Minimizing $\mathcal{L}_{\text{energy}}$ with antiferromagnetic couplings drives the circuit toward the ground state of the learned Hamiltonian,

steering the architecture toward spatially correlated quantum features.

3) *Positional Awareness of the Hamiltonian:* The Hamiltonian is rendered spatially aware through two coupled mechanisms. First, the per-qubit positional rotation θ_{position} depends on the patch-grid coordinate through a learned projection, so distinct spatial positions inject different inputs into the same Hamiltonian topology. Second, because the circuit output state $|\Psi_n\rangle$ depends jointly on the patch feature and the learned positional angle, the effective energy surface experienced by each patch is distinct even when underlying internal features align.

4) *Total Training Objective:* The complete objective function combines pixel-level fidelity with quantum-state regularization:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{rec}} + \lambda \mathcal{L}_{\text{energy}}, \quad (5)$$

evaluated via the Frobenius-norm reconstruction loss:

$$\mathcal{L}_{\text{rec}} = \frac{1}{N} \sum_{n=1}^N \|D(o_n, e_n) - P_n\|_F^2 \quad (6)$$

where D is the classical decoder, P_n is the ground-truth patch, and λ is the energy regularization coefficient.

5) *Symmetry Analysis:* The Heisenberg correlators define three observable sectors:

$$\Xi_{ZZ} = (\langle Z_0 Z_1 \rangle, \dots, \langle Z_4 Z_5 \rangle), \quad \Xi_{XX}, \quad \Xi_{YY} \quad (7)$$

embedded in $\mathbb{R}^{27}$ through the simultaneously measured observables. Two symmetry reductions act sequentially on this manifold:

- 1) **Positional reduction:** The per-qubit positional rotations belong to $\text{SU}(2)$ (acting as a double cover to the spatial $\text{SO}(3)$ Bloch rotations) and break permutation symmetry between qubit sites, reducing continuous symmetry down to the diagonal phase rotations matching the Cartan sector.
- 2) **Geometric reduction:** The 4×4 patch grid carries the dihedral group of the square,

$$D_4 = \langle r, s \mid r^4 = s^2 = (sr)^2 = e \rangle, \quad (8)$$

with eight elements. Under transformation matrices $g \in D_4$, observable vectors transform as $o_{g(i,j)} \approx \Phi_g o_{i,j}$, decomposing the latent space into irreducible representations:

$$\mathcal{V}_{\text{obs}} \approx m_A A_1 \oplus m_B A_2 \oplus m_C B_1 \oplus m_D B_2 \oplus m_E E. \quad (9)$$

The resulting symmetry stack handles the structural degradation systematically:

$$\text{SU}(2)\text{-variant sector} \xrightarrow{\text{positional encoding}} \text{Cartan sector} \xrightarrow{\text{grid geometry}} D_4. \quad (10)$$

This framework preserves the structural symmetries of the patch grid, allowing the pipeline to distinguish orientation-sensitive boundaries.

6) *Symmetric HVK1D Constraint*: We also evaluate a reflection-constrained chain, denoted Symmetric HVK1D, to test whether explicit coupling symmetry can act as a useful structural prior. For a six-qubit chain with $n_{\text{bonds}} = 5$, the unconstrained HVK1D Hamiltonian contains $3n_{\text{bonds}}$ independent couplings. Symmetric HVK1D ties mirror-related bonds:

$$J_{\alpha}^{(i)} = J_{\alpha}^{(n_{\text{bonds}}-1-i)}, \quad \alpha \in \{x, y, z\}. \quad (11)$$

Equivalently, defining the folded bond index

$$\pi(i) = \min\{i, n_{\text{bonds}} - 1 - i\}, \quad (12)$$

the constrained Hamiltonian is

$$H_{\text{sym}} = \sum_{i=0}^{n_{\text{bonds}}-1} \sum_{\alpha \in \{x, y, z\}} J_{\alpha}^{(\pi(i))} \sigma_{\alpha}^{(i)} \sigma_{\alpha}^{(\pi(i))}. \quad (13)$$

This reduces the number of independent coupling parameters from $3n_{\text{bonds}}$ to $3\lceil n_{\text{bonds}}/2 \rceil$ and enforces left-right consistency in the latent spin chain. The constraint is weaker than the full 2D grid prior, but it provides an interpretable test of whether image reconstruction benefits from Hamiltonian parameter tying before adding a full lattice topology.

III. HVK ARCHITECTURE DESIGN

The pipeline consists of an operational flow defined by:

$$\begin{aligned} \text{Image} &\rightarrow \{\text{Patches}\} \rightarrow \text{MPS} \rightarrow \text{46-D Feature} \\ &\rightarrow \text{6-D Vector} \rightarrow \text{Latent} \\ &\rightarrow \text{Decoder} \rightarrow \text{Reconstruction}. \end{aligned} \quad (14)$$

A. Tensor Network Optimization

1) *Patch Preparation*: An input image $\mathbf{I} \in \mathbb{R}^{H \times W}$ is normalized and resized to a canonical resolution of 256×256 pixels. It is partitioned into a regular grid of non-overlapping 64×64 patches:

$$\Pi = \{P_{ij} \in \mathbb{R}^{64 \times 64} : i, j \in \{0, 1, 2, 3\}\}, \quad (15)$$

yielding $N = 16$ patches. Each patch carries a normalized center coordinate $(r_i, c_j) = (i/4, j/4) \in \mathbb{R}^2$.

2) *MPS Feature Extraction*: Each patch P_{ij} is unrolled into a flat vector $\mathbf{v}_n$, ℓ_2 -normalized, and reshaped into a 12-site spin tensor product state:

$$|\psi_n\rangle = \frac{\mathbf{v}_n}{\|\mathbf{v}_n\|_2} \in (\mathbb{C}^2)^{\otimes 12}, \quad (16)$$

satisfying $2^{12} = 4096$ variables. The state is decomposed into a Matrix Product State via sequential Singular Value Decomposition (SVD) with a strict rank truncation set to bond dimension $\chi = 4$:

$$|\psi_n\rangle \approx \sum_{\sigma_1, \dots, \sigma_{12}} A_{\sigma_1} A_{\sigma_2} \dots A_{\sigma_{12}} |\sigma_1 \dots \sigma_{12}\rangle, \quad (17)$$

where $A_{\sigma_i} \in \mathbb{C}^{\chi_{i-1} \times \chi_i}$. This provides a genuine quantum-aligned feature envelope protecting local correlations.

3) *Observable Feature Extraction*: Three distinct measurement families are extracted from the underlying MPS to constitute a 46-dimensional classical feature vector $\mathbf{f}_n$:

- **Local Pauli Expectations**: $\langle Z_i \rangle$ and $\langle X_i \rangle$ for all 12 sites, contributing 24 features.
- **Nearest-Neighbor Correlations**: $\langle Z_i Z_{i+1} \rangle$ across 11 spatial bonds, contributing 11 features.
- **Bipartite Entanglement Entropy**: Calculated across all 11 split configurations via singular values $\lambda_k^{(i)}$:

$$p_k^{(i)} = \frac{|\lambda_k^{(i)}|^2}{\sum_k |\lambda_k^{(i)}|^2}, \quad S_i = - \sum_k p_k^{(i)} \log(p_k^{(i)} + \varepsilon), \quad (18)$$

with $\varepsilon = 10^{-8}$, yielding 11 features.

The final combined configuration produces $\mathbf{f}_n \in \mathbb{R}^{46}$.

4) *Two-Dimensional Fourier Positional Encoding*: Spatial information is injected via a sinusoidal map to break coordinate invariance [17]. For coordinate pair (r_i, c_j) , the vector $e \in \mathbb{R}^8$ is defined by:

$$\begin{aligned} e_{4k} &= \sin(\omega_k \pi r_i), & e_{4k+1} &= \cos(\omega_k \pi r_i), \\ e_{4k+2} &= \sin(\omega_k \pi c_j), & e_{4k+3} &= \cos(\omega_k \pi c_j), \end{aligned} \quad (19)$$

for $k \in \{0, 1\}$ with base frequencies ω_k . This is projected to rotational constraints via a learned map $W_{\text{pos}} \in \mathbb{R}^{6 \times 8}$:

$$\theta_{\text{position}} = W_{\text{pos}} \cdot e \in \mathbb{R}^6. \quad (20)$$

5) *Variational Quantum Circuit*: The raw 46-D feature matrix $\mathbf{f}_n$ is linearly compressed to a 6-D vector $x_n \in \mathbb{R}^6$ prior to circuit injection. For each patch, a 6-qubit register initialized to $|0\rangle^{\otimes 6}$ undergoes three operations:

- 1) **Angle Embedding**: Feature inputs are mapped via Y -axis rotations:

$$U_{\text{enc}}(x_n) = \bigotimes_{i=0}^5 R_Y(x_n^{(i)}), \quad R_{\theta} = e^{-i\theta Y/2}. \quad (21)$$

- 2) **Positional Modulation**: Spatial adjustments are handled via:

$$U_{\text{pos}}(\theta_{\text{pos}}, n) = \bigotimes_{i=0}^5 R_Z(\theta_{\text{position}}^{(i)}). \quad (22)$$

- 3) **Strongly Entangling Layers**: Parameterized operations use $L = 2$ layers containing single-qubit rotations and CNOT rings:

$$U_{\text{var}}(\theta) = \prod_{l=1}^L \left[\text{CNOT-Ring} \cdot \bigotimes_{i=0}^5 \text{Rot}(\phi_i^l, \theta_i^l, \omega_i^l) \right], \quad (23)$$

initialized safely to avoid barren plateaus [1]. The state is output as:

$$|\Psi_n\rangle = U_{\text{var}}(\theta) \cdot U_{\text{pos}}(\theta_{\text{pos}}, n) \cdot U_{\text{enc}}(x_n) |0\rangle^{\otimes 6}. \quad (24)$$

6) *Observable Latent Space*: A total of 27 real expectations are evaluated per execution pass:

$$\Phi_n = (\langle X_i \rangle, \langle Z_i \rangle, \langle Z_i Z_{i+1} \rangle, \langle X_i X_{i+1} \rangle, \langle Y_i Y_{i+1} \rangle) \in [-1, 1]^{27}, \quad (25)$$

representing a structured projection $\Phi : \mathbb{CP}^{63} \rightarrow \mathbb{R}^{27}$.

7) *Classical Decoding and Patch Reconstruction*: The extracted matrix is standardized to standard normal form per column and passed to a classical multi-layer perceptron (MLP) decoder alongside the spatial encodings:

$$\tilde{P}_n = D(\text{Standardize}(\Phi_n), e_n). \quad (26)$$

The architecture utilizes a mirrored transpose structure with expanding channel maps ($4 \rightarrow 8 \rightarrow 16 \rightarrow 3$) coupled with cubic upsampling steps to output the final pixel arrays [20].

IV. METHODOLOGY

A. Datasets

1) *Monalisa Single-Image Benchmark*: The baseline canonical single-image analysis uses a 256×256 grayscale layout of *monalisa.jpg*, divided into 16 evaluation regions of size 64×64 pixels [16].

2) *CIFAR-10 Subset*: To capture structural generalization capabilities, a target subset of five CIFAR-10 images is processed at 32×32 resolution for rapid method comparison. Performance tracking uses per-image and aggregate metric repositories generated by the benchmark scripts [18].

B. Benchmarking Implementations

We execute validation passes using eight distinct comparison configurations:

- **1D HVK**: A control setup utilizing a linear chain of 6 qubits linked via 5 nearest-neighbor bonds to generate the 27-dimensional observable vector.
- **Symmetric HVK1D**: A reflection-constrained 1D chain with mirror-tied Hamiltonian couplings, used to isolate the role of explicit coupling symmetry.
- **2D HVK**: An advanced geometric layout arranging the 6 qubits on a 2×3 lattice to match patch boundary configurations natively, producing a 19-dimensional latent signature.
- **CNN**: A compact convolutional reconstruction baseline optimized directly in image space.
- **MLP**: A direct pixel-space multilayer perceptron baseline.
- **Autoencoder**: A classical bottleneck reconstruction model used as a non-Hamiltonian compression reference.
- **Convolutional GAN Autoencoder**: A standard generative adversarial network processing matching patch blocks [18].
- **Parameterized Hamiltonian Learning (PHL)**: A task-mismatched segmentation reference used to establish positioning parameters rather than structural metrics [22].

C. Training Protocol

Optimization passes utilize the Adam optimizer with learning rates set to $\eta = 0.003$ for the 1D variant and $\eta = 0.004$ for the 2D configuration [18]. Parameter targets are jointly updated over 200 epochs under a rigid loss factor ($\lambda = 0.01$). Initialization parameters scale according to established safe limits [1], [16], [17].

D. Evaluation Metrics

1) *Reconstruction Precision*: The system gauges reconstruction performance using Mean Squared Error (MSE) and Peak Signal-to-Noise Ratio (PSNR):

$$\text{MSE}_A(t) = \frac{1}{N \cdot 64^2} \sum_{n=1}^N \|\tilde{P}_{n,t}^{(A)} - P_n\|_F^2, \quad (27)$$

$$\text{PSNR}_A(t) = 20 \log_{10} \left(\frac{1}{\sqrt{\text{MSE}_A(t)}} \right). \quad (28)$$

2) *Hamiltonian Energy Tracking*: The state optimization depth is directly tracked via the underlying Heisenberg expectation tracking:

$$\langle E \rangle_t = \frac{1}{N} \sum_{n=1}^N \sum_i \left(J_z^{(i)} \langle Z_i Z_{i+1} \rangle_{n,t} + J_x^{(i)} \langle X_i X_{i+1} \rangle_{n,t} + J_y^{(i)} \langle Y_i Y_{i+1} \rangle_{n,t} \right). \quad (29)$$

3) *Order Parameters and Phase Diagnostics*: Longitudinal order parameters track system magnetization states alongside lattice correlation configurations:

$$C_A(t) = \frac{1}{15N} \sum_{n=1}^N \sum_{i=0}^4 \left(\langle Z_i Z_{i+1} \rangle_{n,t} + \langle X_i X_{i+1} \rangle_{n,t} + \langle Y_i Y_{i+1} \rangle_{n,t} \right). \quad (30)$$

Critical structural transition markers are flagged using windowed variance deviations:

$$|C_A(t_c) - \text{median}(C_A)| > 2 \cdot \text{std}(C_A). \quad (31)$$

4) *IBM Heron Hardware Probe Metrics*: For the hardware study, the trained reconstruction loop is not executed end-to-end on the QPU. Instead, compact HVK-style circuits are generated from image patches and submitted to IBM Heron hardware. The measured bitstrings define a hardware-side proxy:

$$\mathcal{L}_{\text{Heron}} = (1 - |\langle M_z \rangle|) + \frac{1}{2} (1 - \langle ZZ \rangle_{\mathcal{E}}), \quad (32)$$

where $|\langle M_z \rangle|$ is the absolute longitudinal order and $\langle ZZ \rangle_{\mathcal{E}}$ is the mean nearest-neighbor or grid-edge ZZ correlation. Lower values indicate stronger ordered and correlated measurements. This quantity is a hardware proxy for latent-state organization, not a pixel-space reconstruction loss.

V. EXPERIMENTAL RESULTS

A. CIFAR-10 Subset Results

Table I compiles the structural reconstruction performance metrics evaluated across the five-image CIFAR-10 subset. Figure 1 visualizes the same aggregate comparison.

The CIFAR subset highlights a nuanced ranking. HVK2D reduces mean MSE by 83.0% relative to the GAN baseline and by 94.6% relative to the autoencoder, while improving over HVK1D by 28.4% in MSE. The CNN and MLP references remain stronger purely classical pixel-space baselines on this small subset; this is expected because they optimize directly in

CIFAR Reconstruction Metric Benchmark Across Architectures




Fig. 1. CIFAR-10 subset reconstruction benchmark across all tested architectures. The left panel compares mean and best reconstruction error on a log scale; the right panel reports mean PSNR with mean MSE annotations.

TABLE I
CIFAR-10 SUBSET EVALUATION (FIVE-IMAGE AVERAGE)

Model	Mean MSE	Mean PSNR	Mean SSIM
Autoencoder	8.604×10^{-3}	20.78	0.8818
CNN	1.103×10^{-4}	40.93	0.9989
GAN	2.732×10^{-3}	26.07	0.9694
HVK1D	6.487×10^{-4}	32.40	0.9931
HVK2D	4.642×10^{-4}	34.48	0.9958
MLP	7.087×10^{-7}	75.13	0.9999
Symmetric HVK1D	6.402×10^{-4}	32.89	0.9939
PHL	4.342×10^{-2}	14.01	0.7879

TABLE II
MONALISA RECONSTRUCTION PERFORMANCE ACROSS METHODS IN THE UNIFIED BASELINE SUITE

Model	Best MSE	Best PSNR (dB)	Best SSIM
Autoencoder	1.2929×10^{-2}	18.88	0.7605
CNN	4.7789×10^{-5}	43.21	0.9993
GAN	2.0129×10^{-3}	26.96	0.9692
HVK1D	7.9668×10^{-4}	30.99	0.9879
HVK2D	5.1973×10^{-4}	32.84	0.9922
MLP	8.0966×10^{-8}	70.92	0.9999
Symmetric HVK1D	1.2059×10^{-3}	29.19	0.9814
PHL [22]	4.3750×10^{-2}	13.59	0.7592

image space without a constrained quantum-observable bottleneck. HVK's contribution is therefore not raw dominance over every classical architecture, but the combination of competitive reconstruction with interpretable Hamiltonian order, energy, and correlation channels.

B. Monalisa Single-Image Results

Cross-architecture reconstruction comparisons for the structural *Monalisa* validation target are summarized in Table II. The table uses the unified baseline suite so that HVK, Symmetric HVK1D, CNN, MLP, autoencoder, GAN, and PHL are compared under the same metric export path.

PHL is included as a Hamiltonian learning reference, but its native output is segmentation-oriented; its additional Dice and IoU metrics are therefore not directly comparable to pixel reconstruction MSE.

C. 1D vs. 2D Topological Ablation Study

Table III highlights the operational adjustments caused by restructuring the lattice configuration from a 1D chain to a 2D adjacent grid.




Fig. 2. Unified Monalisa benchmark metric comparison across HVK variants and classical baselines.

D. MPS Bond-Dimension Sweep Analysis

An optimization sweep across varying tensor sizes targets internal feature generation performance. Setting $\chi = 1$ gives an initial structural encoder MSE of 9.77×10^{-3} (PSNR of 20.10 dB). Scaling up to $\chi = 4$ improves base values to an MSE of 4.44×10^{-4} and a PSNR of 33.52 dB [18], verifying that the subsequent quantum layers supply the extra margin to push performance to the final 40.70 dB mark.

E. IBM Heron Hardware Validation

We executed a compact HVK2D hardware probe on IBM Heron processor `ibm_fez`. A 64×64 Monalisa image was divided into 64 non-overlapping 8×8 patches, and one six-qubit HVK2D circuit was submitted per patch at epoch




Fig. 3. Representative CIFAR frog reconstruction outputs under the 1D chain, reflection-constrained Symmetric HVK1D, and 2D grid topologies. This example is used because the reconstruction differences are visually clearer than the cat sample.




Fig. 4. Order-parameter diagnostics for the same CIFAR frog sample, showing how the Hamiltonian latent trajectory changes under chain, symmetric-chain, and 2D-grid constraints.




Fig. 5. Monalisa Hamiltonian order-parameter tracking for the 1D, Symmetric HVK1D, and 2D variants. These plots report graphical order diagnostics rather than original-versus-reconstructed image panels.




Fig. 6. Bond-dimension sweep showing how increasing the MPS truncation rank changes the pre-quantum feature quality for 1D and 2D HVK configurations.

TABLE III
LATTICE TOPOLOGY PERFORMANCE ADJUSTMENTS

Metric	1D HVK	2D HVK
Critical Epoch	11	2
Final Energy	-0.1154	-0.2411
Order Parameter $\langle M_z \rangle$	0.0411	0.0094
Maximum Susceptibility $\chi_{\max}$	1.15×10^{-4}	6.24×10^{-4}
Final PSNR (dB)	34.72	40.70

TABLE IV
MONALISA ORDER-PARAMETER TRACKING AT EPOCH 199

Model	Final Energy	Final $\langle M_z \rangle$	Final $ \langle M_z \rangle $	Peak Susceptibility Epoch
HVK1D	-4.4877	0.2149	0.2225	3
HVK2D	-6.0070	0.3535	0.3674	138
Symmetric HVK1D	-3.8677	-2.14×10^{-4}	0.0358	3

TABLE V
IBM HERON HVK2D PATCH-PROXY RESULTS

Statistic	Value
Backend	ibm_fez
QPU family	Heron r2
Patch circuits	64
Shots per circuit	100
Mean proxy loss	0.4353
Proxy loss std.	0.0473
Minimum proxy loss	0.3043
Maximum proxy loss	0.5543
Mean $ \langle M_z \rangle $	0.7778
Mean $\langle ZZ \rangle$	0.5738

label 200 with 100 shots per circuit. The aggregate hardware statistics are reported in Table V.

The hardware results show a stable positive ordered re-

sponse. The mean absolute longitudinal order is 0.7778, and the mean ZZ correlation is 0.5738, indicating that the QPU-measured states retain strong alignment and nearest-neighbor structure across the patch set. The patch-proxy map in Fig. 7 visualizes the spatial variation of this response. It should not be interpreted as a full reconstructed image; rather, it is a real-

hardware map of Hamiltonian latent organization induced by image patches.




Fig. 7. Original Monalisa image beside the IBM Heron HVK2D hardware patch-proxy map at epoch label 200.




Fig. 8. Initial IBM Heron hardware probe metrics comparing HVK1D and HVK2D patch circuits.

F. Empirical Observations

The experimental results support four principal observations. First, the reconstruction error decreases substantially once MPS-derived patch descriptors are coupled to the Hamiltonian latent space. This indicates that the MPS layer is not merely a preprocessing compression step; rather, it supplies a structured feature basis on which the variational circuit can form physically interpretable correlation channels. The bond-dimension sweep further shows that increasing the MPS rank improves the quality of the pre-quantum representation, but the final HVK reconstruction still benefits from the subsequent Hamiltonian and decoder stages.

Second, the 2D HVK exhibits stronger reconstruction performance than the 1D chain model on the *Monalisa* experiment. The final MSE decreases from 3.3754×10^{-4} for the 1D configuration to 8.5114×10^{-5} for the 2D configuration, while PSNR increases from 34.72 dB to 40.70 dB. This behavior is consistent with the hypothesis that matching the circuit interaction graph to the spatial adjacency of image patches improves the inductive bias of the reconstruction pipeline.

Third, the Hamiltonian energy and order-parameter diagnostics provide information that is not available from purely image-space baselines. The 2D HVK reaches a lower final energy and displays an earlier, larger susceptibility peak than the 1D model. Since the system is finite and trained by gradient descent, this peak should not be interpreted as a thermodynamic phase transition. It is more appropriately read as a finite-

size reorganization of the learned observable manifold during optimization.

Fourth, the GAN, CNN, MLP, and PHL baselines clarify the task-specific role of HVK. CNN and MLP models optimize directly in pixel space and can be highly effective on small low-resolution subsets. The GAN baseline optimizes image-space realism but does not expose Hamiltonian energy or spin-order diagnostics. The PHL baseline is structurally Hamiltonian, but its output is a segmentation mask rather than a reconstructed image. HVK occupies a different position: it uses Hamiltonian structure as a latent regularizer inside a reconstruction model, thereby producing both pixel-level outputs and interpretable quantum-observable trajectories.

Fifth, the IBM Heron run confirms that the compact HVK2D circuits are executable on present-day superconducting quantum hardware. The hardware proxy shows a non-trivial spread across all 64 patches, with proxy loss ranging from 0.3043 to 0.5543. The weak direct correlation between hardware-proxy brightness and original patch intensity indicates that this signal is not simply grayscale brightness; it is better understood as a QPU-measured correlation response. Connecting this response to a trained decoder is the next step toward full hardware-in-the-loop reconstruction.

VI. DISCUSSION

A. Reconstruction Error Analysis

HVK2D's 83.0% CIFAR MSE reduction relative to the GAN demonstrates that adversarial objectives can introduce structural instabilities during high-fidelity generation. In the *Monalisa* case, moving from the 1D chain to the 2D lattice reduces final MSE from 3.3754×10^{-4} to 8.5114×10^{-5} . HVK enforces structural stability through direct pixel losses paired with smooth physical regularizers derived from learned Hamiltonian states.

B. Architectural Differentiation Matrix

Table VI contrasts the structural properties of HVK against current reference frameworks.

C. The Topology-Alignment Principle

The +5.98 dB performance advantage of the 2D HVK configuration over its 1D counterpart validates the *topology-alignment principle*: aligning the qubit layout of the variational circuit with the spatial adjacency of the image patches prevents conflicting optimization trajectories. This reinforces structural coordination across patch boundaries along the grid edges.

D. Quantitative Channel Economy

The application of 6 qubits maps precisely to 27 or 19 target measurement dimensions, tracking the specific number of local and nearest-neighbor interactions. HVK maintains compact scaling rules while outpacing over-parameterized convolutional models.

TABLE VI
ARCHITECTURAL DIFFERENTIATION MATRIX

Axis	Classical AE [20]	GAN AE [18]	Standard QAE [13]	HVK (Ours)
Latent Space	Continuous $\mathbb{R}^d$	Continuous $\mathbb{R}^d$	Quantum Register	Pauli Correlator Vector
Spatial Bias	Implicit Convolutions	Implicit Convolutions	None	2D Fourier + D_4 Symmetries
Regularizer	Bottleneck Size	Discriminator Loss	Swap-Test Fidelity	Hamiltonian Energy $\mathcal{L}_{\text{energy}}$
Group Alignment	None	None	Reference Swaps	SO(3) Cartan + D_4 Grid
Target Task	Reconstruction	Reconstruction	Compression	Spatially Aware Reconstruction

E. Interpretive Limits

We explicitly acknowledge the following bounds on our empirical claims:

- The CIFAR-10 evaluations are currently bounded to a five-image validation subset.
- CNN and MLP baselines remain stronger direct pixel-space baselines on the small CIFAR subset.
- The IBM Heron study measures a hardware proxy based on order and ZZ correlation; it is not a full QPU-executed decoder reconstruction.
- Larger hardware studies require either batching, error mitigation, or selected epoch snapshots rather than full 200-epoch QPU training.

VII. CONCLUSION

We introduced HVK, a Hamiltonian-guided tensor-network framework for spatially aware quantum image reconstruction. HVK combines an MPS encoder, 2D Fourier positional inputs, and a parameterized variational circuit under a data-adaptive Heisenberg energy regularizer.

On the CIFAR-10 subset, HVK2D achieves a mean PSNR of 34.48 dB and reduces mean squared error by 83.0% relative to the GAN baseline. On the *Monalisa* target, the 2D HVK setup delivers a PSNR of 40.70 dB, exceeding the 1D control by 5.98 dB and the GAN by 10.78 dB. Quantum state diagnostics support the topology-alignment principle, confirming that structured geometric boundaries accelerate convergence and yield deeper energy states. The IBM Heron study further demonstrates that compact HVK circuits can be executed on real Heron hardware, producing stable order and ZZ correlation measurements across a full 64-patch *Monalisa* probe.

Future efforts will focus on running full-scale CIFAR-10 dataset evaluations with the 2D lattice structure, integrating Heron measurements with the trained decoder for true hardware-in-the-loop reconstruction, implementing error mitigation, and assessing performance against a broader collection of conditional generative baselines.

REFERENCES

- [1] M. Cerezo et al., “Variational quantum algorithms,” *Nature Reviews Physics*, vol. 3, no. 9, pp. 625–644, 2021.
- [2] J. Tilly et al., “The Variational Quantum Eigensolver: A review of methods and best practices,” *Physics Reports*, vol. 986, pp. 1–128, 2022.
- [3] B. Bauer, S. Bravyi, M. Motta, and G. K.-L. Chan, “Quantum Algorithms for Quantum Chemistry and Quantum Materials Science,” *Chemical Reviews*, vol. 120, no. 22, pp. 12685–12717, 2020.
- [4] S. Chakraborty, S. B. Mandal, and S. H. Shaikh, “Quantum image processing: challenges and future research issues,” *International Journal of Information Technology*, vol. 14, no. 1, pp. 475–489, 2018.
- [5] R. Zhou, W. Hu, P. Fan, and H. Ian, “Quantum realization of the bilinear interpolation method for NEQR,” *Scientific Reports*, vol. 7, no. 1, p. 2511, 2017.
- [6] X. Yao et al., “Quantum Image Processing and Its Application to Edge Detection: Theory and Experiment,” *Physical Review X*, vol. 7, no. 3, p. 031041, 2017.
- [7] J. I. Cirac, D. Perez-Garcia, N. Schuch, and F. Verstraete, “Matrix product states and projected entangled pair states: Concepts, symmetries, theorems,” *Reviews of Modern Physics*, vol. 93, no. 4, p. 045003, 2021.
- [8] J. I. Latorre, “Image compression and entanglement,” arXiv:quant-ph/0510031, 2005.
- [9] Z. Han, J. Wang, H. Fan, L. Wang, and P. Zhang, “Unsupervised Generative Modeling Using Matrix Product States,” *Physical Review X*, vol. 8, no. 3, p. 031012, 2018.
- [10] S. M. Fei, “Compressed Sensing Based on Tensor Network Machine Learning,” *Physical Science & Biophysics Journal*, vol. 5, no. 1, p. 16000174, 2021.
- [11] D. Guala et al., “Practical overview of image classification with tensor-network quantum circuits,” *Scientific Reports*, vol. 13, no. 1, p. 4427, 2023.
- [12] A. Haga, “Quantum annealing-based computed tomography using variational approach for a real-number image reconstruction,” *Physics in Medicine and Biology*, vol. 69, no. 4, p. 02155, 2024.
- [13] X. Zhai, T. Xiao, J. Huang, J. Fan, and G. Zeng, “Quantum neural compressive sensing for ghost imaging,” *Physical Review Applied*, vol. 23, no. 1, p. 014018, 2025.
- [14] M. A. Nau, A. H. Vija, W. Gohn, M. P. Reymann, and A. Maier, “Exploring the Limitations of Hybrid Adiabatic Quantum Computing for Emission Tomography Reconstruction,” *Journal of Imaging*, vol. 9, no. 10, p. 221, 2023.
- [15] B. Craps, M. D. Clerck, O. Evnin, and P. Hacker, “Integrability and complexity in quantum spin chains,” *SciPost Physics*, vol. 16, no. 2, p. 041, 2024.
- [16] “Mathematical Framework for HVK,” internal technical report, June 2026.
- [17] R. Xu, X. Wang, K. Chen, B. Zhou, and C. C. Loy, “Positional Encoding as Spatial Inductive Bias in GANs,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021, pp. 13564–13573.
- [18] “Methodology and Experimental Analysis of the Hamiltonian Vision Kernel,” internal laboratory reference, 2026.
- [19] J. Y. Araz and M. Spannowsky, “Classical versus quantum: Comparing tensor-network-based quantum circuits on Large Hadron Collider data,” *Physical Review A*, vol. 106, no. 6, p. 062423, 2022.
- [20] D. Slabbert and F. Petruccione, “Hybrid Quantum-Classical Feature Extraction approach for Image Classification using Autoencoders and Quantum SVMs,” arXiv:2410.18814, 2024.
- [21] Z. Liu, P.-X. Shen, W. Li, L.-M. Duan, and D.-L. Deng, “Quantum capsule networks,” *Quantum Science and Technology*, vol. 8, no. 1, p. 015016, 2022.
- [22] J. Shi, W. Wang, X. Lou, S. Zhang, and X. Li, “Parameterized Hamiltonian Learning With Quantum Circuit,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 45, no. 5, pp. 6086–6095, 2022.